
Appendix B

Solutions

Solutions for Practice 2: Installing the Oracle Database Software

Background: In the practices of this course, you assume the role of a database administrator (DBA). The operating system (OS) accounts on your computer are:

- The `oracle` user with a password of `oracle`
- The `root` user with a password of `oracle`

The system administrator has set up the OS so that it is ready for the installation, and the installation media is staged at `/stage/Disk1`. Perform the following tasks as the default `oracle` OS user, unless otherwise indicated.

After installing the software, execute the `lab_02_05.sh` script, which creates a listener for you.

Note: Completing this practice is critical for all following practice sessions.

1. Install the Oracle database software as the `oracle` user. Navigate to the `/stage/Disk1` directory, and start the Oracle Universal Installer (OUI) by entering `./runInstaller`.
 - a) Right-click your desktop and select **Open Terminal**, and then enter:

```
$ cd /stage/Disk1
$ ./runInstaller
```

2. Select your installation method for OUI.
 - a) On the Installation Method page, select **Basic Installation**, and confirm the following settings:


Object	Setting
Database Home Location	<code>/u01/app/oracle/product/10.2.0/db_1</code>
Installation Type	Enterprise Edition
UNIX DBA Group	<code>oinstall</code>
Create Starter Database	<i>Deselected</i>

Note: Ensure that you deselect the Create Starter Database option.

- b) Click Next.
- c) You are now on the page titled “Specify Inventory directory and credentials.” Accept `/u01/app/oracle/oraInventory` as **inventory directory** and `oinstall` as **Operating System group name**. Click Next.

Solutions for Practice 2: Installing the Oracle Database Software (continued)

OUI is loading the products list and checking product-specific prerequisites.




- d) After OUI has finished its prerequisite checks on the Product-Specific Prerequisite Checks page, click **Next**.
3. When the prerequisite checks are finished, the Summary page is displayed.
 - a) Click **Install** to begin your installation.

Estimated installation time is 10–15 minutes. However, varying environments can greatly influence this estimate.

4. When the “Execute Configuration scripts” page appears, follow the instructions on that page, accept the default for the local bin directory, and then finish your installation with OUI.

Solutions for Practice 2: Installing the Oracle Database Software (continued)



- a) Right-click your desktop and select **Open Terminal**.
- b) Run the `orainstRoot.sh` and `root.sh` scripts as root, and then click **Continue**.

```
$ su
# password: oracle <root password, does not appear on the screen>
# cd /u01/app/oracle/oraInventory
# ./orainstRoot.sh
# cd /u01/app/oracle/product/10.2.0/db_1
# ./root.sh
```

- c) Accept the default for the local bin directory.

Solutions for Practice 2: Installing the Oracle Database Software (continued)

```
[oracle@EDRSR4P1 oracle]$ su
Password:
[root@EDRSR4P1 oracle]# cd /u01/app/oracle/oraInventory
[root@EDRSR4P1 oraInventory]# ./orainstRoot.sh
Changing permissions of /u01/app/oracle/oraInventory to 770.
Changing groupname of /u01/app/oracle/oraInventory to oinstall.
The execution of the script is complete
[root@EDRSR4P1 oraInventory]# cd /u01/app/oracle/product/10.2.0/db_1
[root@EDRSR4P1 db_1]# ./root.sh
Running Oracle10 root.sh script...

The following environment variables are set as:
    ORACLE_OWNER= oracle
    ORACLE_HOME=  /u01/app/oracle/product/10.2.0/db_1


Enter the full pathname of the local bin directory: [/usr/local/bin]:
Copying dbhome to /usr/local/bin ...
Copying oraenv to /usr/local/bin ...
Copying coraenv to /usr/local/bin ...

Creating /etc/oratab file...
Entries will be added to the /etc/oratab file as needed by
Database Configuration Assistant when a database is created
Finished running generic part of root.sh script.
Now product-specific root actions will be performed.
[root@EDRSR4P1 db_1]# █
```

- d) Enter **exit** to exit the root OS user.
- e) Close the terminal window, and then click **OK** on the “Execute Configuration scripts” page.

The End of Installation page appears.

Solutions for Practice 2: Installing the Oracle Database Software (continued)



f) Make a note of your URLs. You will use them in later practice sessions.

iSQL*Plus URL: _____

iSQL*Plus DBA URL: _____

g) Optionally, click **Installed Products**, review the product inventory, then click **Close**.

h) Click **Exit**, and then click **Yes** to leave OUI.

5. Execute the `lab_02_05.sh` script, which is in the `/home/oracle/labs` directory. This script uses the `netca` utility to create a listener for you. If you create your first database with OUI (as part of your installation), then OUI invokes the network configuration assistant, which creates the first listener for you.

a) In a terminal window, enter:

```
cd /home/oracle/labs
./lab_02_05.sh
```

You see an activity log, which should end with a success message. If not, resolve any errors that might have occurred.

Solutions for Practice 3: Creating an Oracle Database

Background: You are about to begin creating your first Oracle database. You anticipate that several similar databases will be needed in the near future. Therefore, you decide to create your ORCL database, as well as a database template and the database creation scripts. Locate the scripts in the `/home/oracle/labs` directory (which is the directory that you use most often throughout this course).

After you create the ORCL database, you execute the `lab_03_03.sh` script, which creates a listener for you.

Note: Completing the database creation is critical for all following practice sessions.

1. Start the Database Configuration Assistant (DBCA).
 - a) Open a terminal window as the `oracle` user; that is, right-click your desktop and select **Open Terminal**.
 - b) To start the DBCA, enter:

```
$ dbca
```

2. Begin the ORCL database creation. Use the General Purpose database template.
 - a) In the DBCA, click **Next** on the Welcome page.
 - b) On the Operations page, select **Create a Database**, and then click **Next**.
 - c) On the Database Templates page, select **General Purpose**, and then click **Show Details**.
 - d) Review the template's details and answer the following questions.

Question 1: How many control files are created?

Answer: 3

Question 2: Would it maximize database availability to multiplex them?

Answer: Yes. (*This will be done in a later practice.*)

Question 3: How many redo log groups are created?

Answer: 3

Question 4: Would it maximize database availability to mirror them?

It depends: No, not in class, because there are already three groups and you have only one physical storage device; but yes, if you can put each group on a different physical storage device.

Solutions for Practice 3: Creating an Oracle Database (continued)

Question 5: What is the database block size (db_block_size)?

Answer: 8 KB

Question 6: What is the value of Sample Schemas?

Answer: Sample Schemas is set to False.

Note: You will change this setting later in this practice to create the HR sample schema.

Question 7: What is the template default for the Database Character Set?

Answer: WE8ISO8859P1

Note: You will change this setting later in this practice to use a Unicode database character set.

- e) Click **Close** to close the Template Details window.
 - f) Click **Next**.
3. Create the ORCL database, as well as the ORCL template and the database generation scripts.
- a) On the Database Identification page, enter **orcl.oracle.com** as Global Database Name. The SID defaults to the database name **orcl**. Click **Next**.
 - b) On the Management Options page, ensure that the following items are selected:
 - Configure the Database with Enterprise Manager
 - Use Database Control for Database Management
 - c) Click **Next**.
 - d) On the Database Credentials page, select **Use the Same Password for All Accounts** and enter **oracle** as Password and Confirm Password. Then, click **Next**.
 - e) On the Storage Options page, select **File System**, and then click **Next**.
 - f) On the Database File Locations page, select **Use Oracle-Managed Files**. Accept the default Database Area, and then click **Next**.
 - g) On the Recovery Configuration page, select **Specify Flash Recovery Area**, and then click **Next**.
 - h) On the Database Content page, select **Sample Schemas**, and then click **Next**.
 - i) On the Memory tabbed page of the Initialization Parameters page, select **Custom** and then select **Automatic** for the **Shared Memory Management** setting.

Solutions for Practice 3: Creating an Oracle Database (continued)

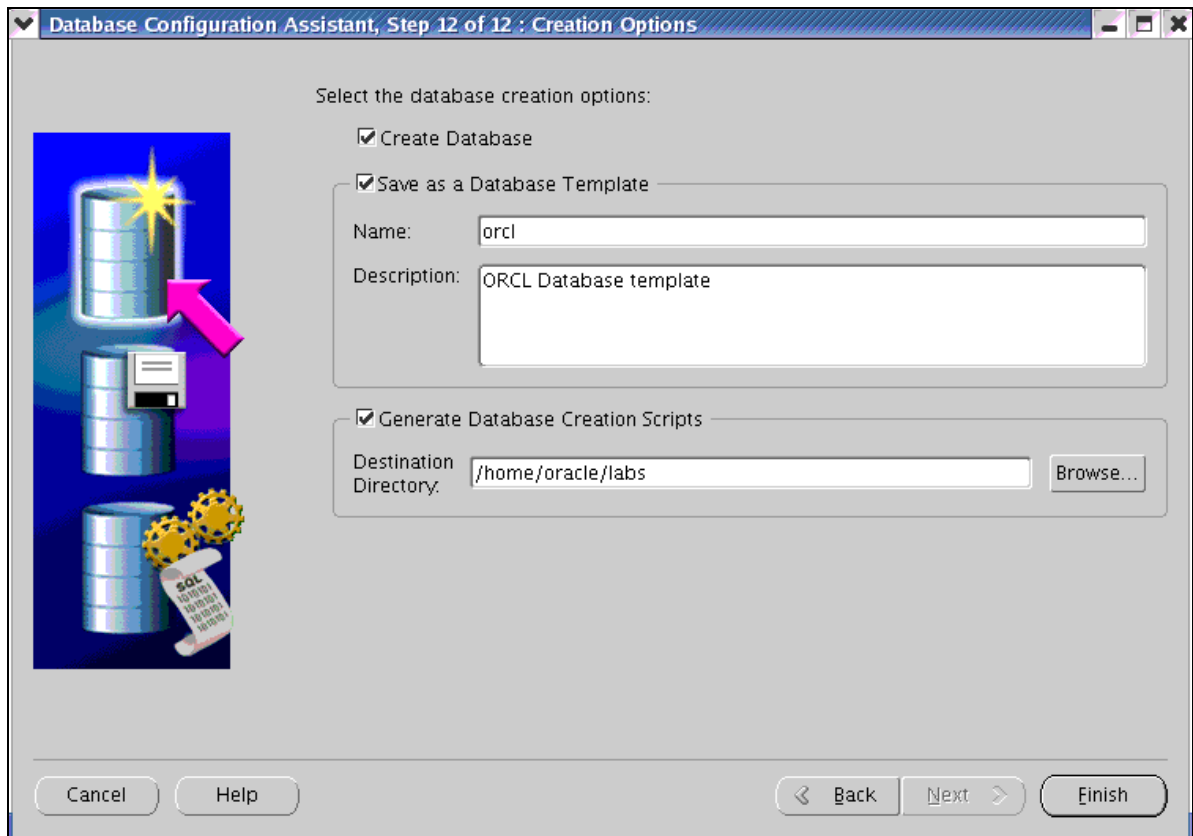
- j) On the Initialization Parameters Character Sets tabbed, select **Use Unicode (AL32UTF8)**.

The screenshot shows the 'Character Sets' tab in the Oracle Database Initialization Parameters window. The 'Database Character Set' section has three radio buttons: 'Use the default' (unselected), 'Use Unicode (AL32UTF8)' (selected), and 'Choose from the list of character sets' (unselected). Below the radio buttons, a text box explains that the default character set is based on the operating system's language setting (WE8ISO8859P1). A dashed line separates the 'Use Unicode' option from the 'Choose from the list' option. Below the dashed line, a text box explains that setting the character set to Unicode (AL32UTF8) enables storing multiple language groups. Below the 'Choose from the list' option, a dropdown menu for 'Database Character Set' is shown, currently displaying 'WE8ISO8859P1 - ISO 8859-1 West European'. Below the 'Database Character Set' section, there are three more dropdown menus: 'National Character Set' (displaying 'AL16UTF16 - Unicode UTF-16 Universal character set'), 'Default Language' (displaying 'American'), and 'Default Date Format' (displaying 'United States').

- k) Review the Sizing and Connection Mode tabbed pages, but do not change any values. Then, click **Next**.
- l) On the Database Storage page, review your file names and locations. Then, click **Next**.

Solutions for Practice 3: Creating an Oracle Database (continued)


- m) On the Creation Options page, select **Create Database**.



- n) Optionally, select all creation options and enter `orcl` as **Name** for the database template, `ORCL Database template` as **Description**, and `/home/oracle/labs` as **Destination Directory**. Then, click **Finish**.
- o) The Confirmation page appears. Review options and parameters, such as `Sample Schemas (true)`, `db_block_size (8KB)`, `sga_target (270MB)`, `undo_management (AUTO)`, `Database Character Set (AL32UTF8)`, and then click **OK**.
- p) Click **OK** to acknowledge that the template has been created. Then, acknowledge the generation of the database scripts (if you selected those options).

Solutions for Practice 3: Creating an Oracle Database (continued)

- q) The DBCA displays the progress of the various installation steps. When the database itself has been created, the DBCA displays essential information about the database. Make note of this information. The Database Control URL will be used in several of the following practice sessions.



- r) Click the **Password Management** button.
- s) Scroll down the Password Management page until you see the HR User Name.

Solutions for Practice 3: Creating an Oracle Database (continued)

- t) Deselect **Lock Account?** and enter **hr** as New Password and Confirm Password. Then, click **OK**.

User Name	Lock Account?	New Password	Confirm Password
SI_INFORMTN_SCHEMA	✓		
OLAPSYS	✓		
MDDATA	✓		
IX	✓		
SH	✓		
DIP	✓		
OE	✓		
SCOTT	✓		
HR		**	**
TSMSYS	✓		
BI	✓		
PM	✓		

- u) Click **Exit** to close the DBCA.

You completed your task to create a database and (optionally) a database template and database generation scripts.


Solutions for Practice 4: Managing the Oracle Instance

Background: You have just installed the Oracle software and created a database. You want to ensure that you can start and stop the database and see the application data.

1. Invoke Enterprise Manager, and log in as the SYS user. Which port number does this database use? You noted this in Practice 3.

Answer: 1158

- a) Open a Web browser.
- b) If using Mozilla, the Select User Profile dialog box appears.



Select **oracle** as the profile, and select **Don't ask at startup**, and then click **Start Mozilla**.

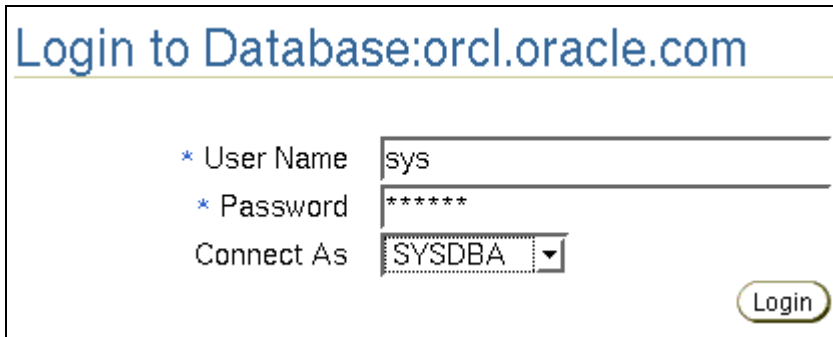
- c) Enter the URL that you wrote down in Practice 3. It has the following format:

`http://hostname:portnumber/em`

The Oracle Enterprise Manager window appears.

Solutions for Practice 4: Managing the Oracle Instance (continued)

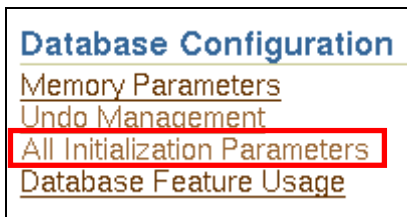
- d) Log in by entering `sys` as **User Name** and `oracle` as **Password**, and selecting `SYSDBA` as **Connect As**. Then, click **Login**.



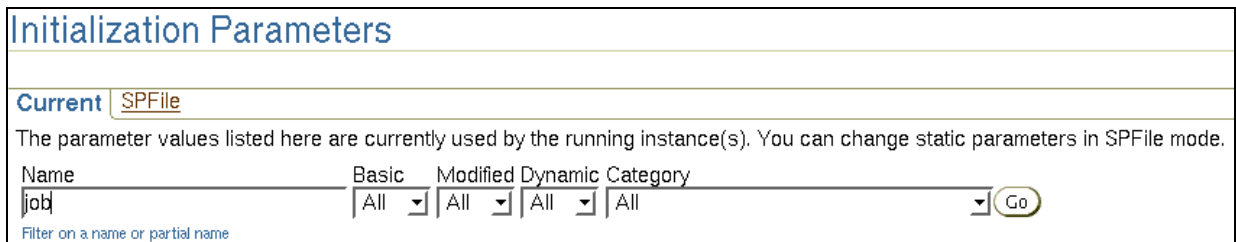
- e) You may receive a security warning. In class, deselect **Alert me whenever I submit information that's not encrypted.**, and then click **Continue**.

The first time you log in after installing the software, the Oracle Database 10g Licensing Information page appears.

- f) To acknowledge this information, click **I agree** at the bottom of the page.
2. View the initialization parameters and set the `JOB_QUEUE_PROCESSES` parameter to 15. What SQL statement is run to do this?
- a) Select **Administration** > Database Administration > Database Configuration > **All Initialization Parameters**.



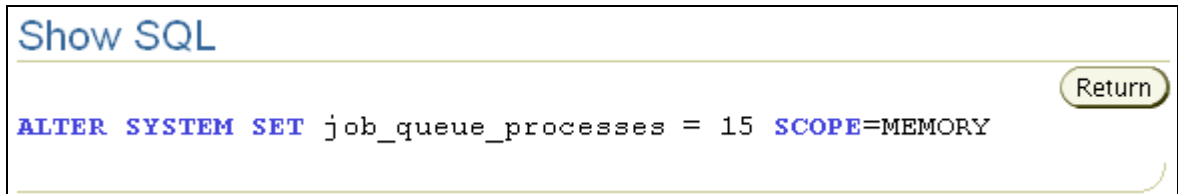
- b) Enter `job` in the **Name** field, and then click **Go**.



- c) When the `JOB_QUEUE_PROCESSES` initialization parameter appears, change its value from 10 to 15.

Solutions for Practice 4: Managing the Oracle Instance (continued)

- d) Click **Show SQL** and note the SQL statement that is going to be run.



Show SQL

ALTER SYSTEM SET job_queue_processes = 15 **SCOPE=MEMORY**

Return

- e) Click **Return**, and then click **Apply**.

3. *Question:* What is the significance of a check in the Dynamic column?

Answer: A “dynamic” parameter can be modified while the database is active.

4. Shut down the database instance by using Enterprise Manager.

Question: What SQL is executed to do this?

- a) In the Enterprise Manager browser session, click the **Database** tab.
- b) Click **Shutdown**.
- c) For **Host Credentials**, enter `oracle` as **Username** and `oracle` as **Password**.
- d) Click **OK**.

The Startup/Shutdown: Confirmation page appears.

- e) Click **Advanced Options** to see the mode for shutting down, but do not change the mode; it should remain as “Immediate.”
- f) Click **Cancel** to return to the previous page.
- g) Click **Show SQL** to view the SQL that is going to be executed to perform the shutdown operation.

Question: What SQL is executed to do this?

Answer: SHUTDOWN IMMEDIATE



Show SQL

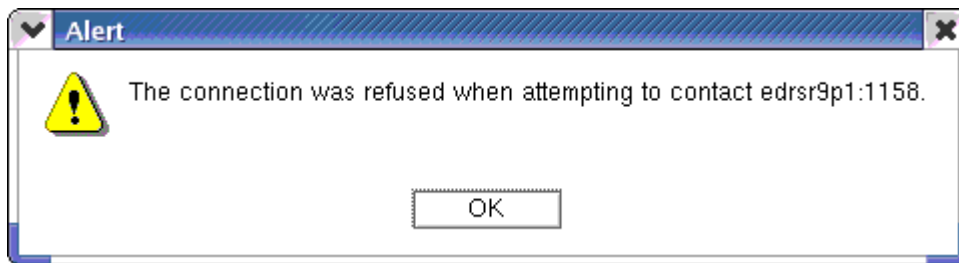
SHUTDOWN immediate

Return

- h) Click **Return**.

Solutions for Practice 4: Managing the Oracle Instance (continued)

- i) Click **Yes** to confirm the shutdown operation.
- j) Click Refresh. If you see the following error during the refresh, click **OK** and continue to refresh. The error will resolve itself.



- k) Note that the **Status** of the instance is now "Down."



- 5. Using SQL*Plus, verify that you are not able to connect as the HR user to a database that has been shut down.
 - a) In the Linux command window, enter the following to attempt to log in to the database:

```
$ sqlplus hr
```
 - b) Enter hr for the password.
 - c) Note the "ORACLE not available" error message.
 - d) Press [Ctrl], [D] to exit the username prompt.
- 6. Use Enterprise Manager to restart the database instance, and then log in as the SYS user again.

Question: What SQL is run to accomplish the database startup?


- a) In Enterprise Manager, click the **Startup** button.
- b) Enter oracle for both **Username** and **Password** in the **Host Credentials** region.
- c) Click **OK**.

Solutions for Practice 4: Managing the Oracle Instance (continued)

- d) The Startup/Shutdown: Confirmation page appears. Click **Show SQL** to view the SQL that is about to run.

Question: What SQL is run to accomplish the database startup?

Answer: STARTUP



- e) Click **Return**.
 - f) Click **Yes** to confirm the startup operation.
 - g) The Startup/Shutdown: Activity Information page appears. Wait for the login page to appear, at which time you can log in as SYS user with the `oracle` password and the SYSDBA privilege.
7. In the alert log, view the phases that the database went through during startup. What are they?
- a) Select **Database** > Related Links > **Alert Log Content**.

Solutions for Practice 4: Managing the Oracle Instance (continued)

- b) Scroll toward the bottom of the log and review the phases of the database during startup. Your alert log may look different from this screenshot, based on different system activities.

```
Fri May 20 09:47:26 2005
Database mounted in Exclusive Mode
Completed: ALTER DATABASE MOUNT
Fri May 20 09:47:26 2005
ALTER DATABASE OPEN
Fri May 20 09:47:26 2005
Thread 1 opened at log sequence 209
  Current log# 1 seq# 209 mem# 0: /u01/app/oracle/oradata/orcl/redo01.log
Successful open of redo thread 1
Fri May 20 09:47:26 2005
MTTR advisory is disabled because FAST_START_MTTR_TARGET is not set
Fri May 20 09:47:26 2005
SMON: enabling cache recovery
Fri May 20 09:47:27 2005
Successfully onlined Undo Tablespace 1.
Fri May 20 09:47:27 2005
SMON: enabling tx recovery
Fri May 20 09:47:27 2005
Database Characterset is AL32UTF8
replication_dependency_tracking turned off (no async multimaster replication found)
Starting background process QMNC
QMNC started with Fri May 20 09:47:30 2005
db_recovery_file_dest_size of 2048 MB is 0.00% used. This is a
user-specified limit on the amount of space that will be used by this
database for recovery-related files, and does not reflect the amount of
space available in the underlying filesystem or ASM diskgroup.
Fri May 20 09:47:33 2005
Completed: ALTER DATABASE OPEN
```

- c) Note that the modes the database goes through during startup are MOUNT and OPEN. Click **OK** to close the alert log.
8. Test access to *iSQL*Plus* for your HR application developers. (Navigation aid: Database > *iSQL*Plus*). Use the Normal role, hr username and password, and the default setting as Connect Identifier. If there is an error accessing *iSQL*Plus*, then start the isqlplus process using the `isqlplusctl start` command at the OS prompt, and then reattempt. After connecting, select the contents of the EMPLOYEES table.

- a) Select **Database > Related Links > iSQL*Plus**.

The *iSQL*Plus* Connection Role page appears.

- b) Notice that the SYSOPER and SYSDBA roles require special setup and authentication for security reasons. Select **Normal**, and then click **Continue**.

Solutions for Practice 4: Managing the Oracle Instance (continued)

- c) If you see an error message saying that the connection was refused, that means you need to start the supporting process. To do this, enter the `isqlplusctl start` command at the OS prompt, and then retry step (b).

```
$ isqlplusctl start
iSQL*Plus 10.2.0.1.0
Copyright (c) 2003, 2005, Oracle. All rights reserved.
Starting iSQL*Plus ...
iSQL*Plus started.
```

- d) On the Login page, enter `hr` as **Username** and **Password** and leave the **Connect Identifier** set to its default value. Click the **Login** button.

The Confirm window appears, with the Password Manager offering to remember this login for you. Click **No**.

- e) Enter the following statement in the workspace, and then click the **Execute** button:

```
SELECT * FROM EMPLOYEES;
```

The content of the `EMPLOYEES` table appears in a tabular form.

Connected as **HR@orcl**

Workspace

Enter SQL, PL/SQL and SQL*Plus statements. Clear

```
SELECT * FROM EMPLOYEES;
```

Execute Load Script Save Script Cancel

EMPLOYEE_ID	FIRST_NAME	LAST_NAME	EMAIL	PHONE_NUMBER	HIRE_DATE	JOB_ID	SALARY
198	Donald	OConnell	DOCONNEL	650.507.9833	21-JUN-99	SH_CLERK	2600
199	Douglas	Grant	DGRANT	650.507.9844	13-JAN-00	SH_CLERK	2600
200	Jennifer	Whalen	JWHALEN	515.123.4444	17-SEP-87	AD_ASST	4400
201	Michael	Hartstein	MHARTSTE	515.123.5555	17-FEB-96	MK_MAN	13000

- f) When you have finished reviewing the information, click **Logout** in the top-right corner of the page.



- g) Click X in the top-right window frame to close the `iSQL*Plus` window.

Solutions for Practice 5: Managing Database Storage Structures

Background: You need to create a new tablespace for the INVENTORY application. You also need to create a database user that is not as privileged as the SYS user.

1. Enter `./lab_05_01.sh` to run a script that creates the DBA1 user. It is located at `/home/oracle/labs`. The password for DBA1 is `oracle`.

- a) Start a Linux command shell by right-clicking your desktop and selecting **Open Terminal**.

- b) Change the current directory to the `labs` directory by entering:

```
$ cd labs
```

- c) Enter the following command to run the script that creates the DBA1 user:

```
$ ./lab_05_01.sh
```

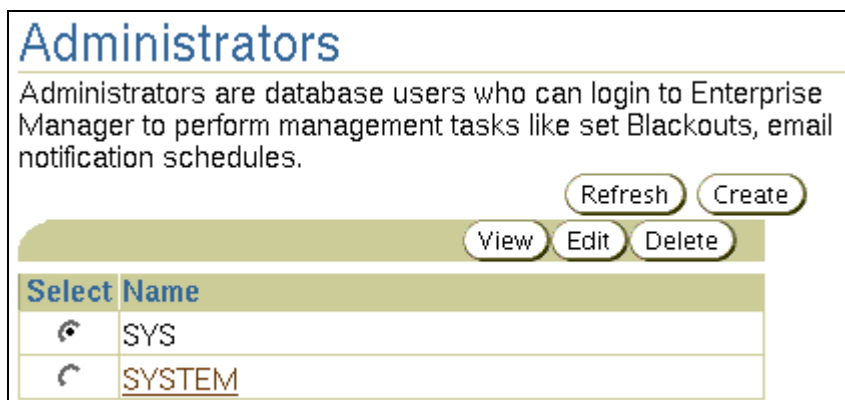
- d) Leave the command shell window open. You will use it again later.

2. Use the Setup link in the top-right corner of Enterprise Manager (EM) to define the DBA1 user as one who can perform administrative tasks in EM. When the non-SYS user is configured, log out as the SYS user and log in as the DBA1 user. Use the DBA1 user to perform the rest of these tasks, unless otherwise indicated.

- a) In the far top-right corner of the EM window, click **Setup**.



- b) Click **Create** to add the DBA1 user to the Administrators list. This will enable the DBA1 user to perform management tasks by using Enterprise Manager.



Solutions for Practice 5: Managing Database Storage Structures (continued)

- c) Enter `dba1` as **Name**, and `oracle` as **Password** and **Confirm Password**. Leave **Email Address** blank, and leave **Super Administrator** selected, and then click **Finish**.
- d) On the Create Administrator: Review page, click **Finish** again.
- e) Now that there is a non-SYS user, click **Logout** in the top-right corner, and then click **Login**.
- f) Enter `DBA1` as **User Name** and `ORACLE` as **Password**, and select `SYSDBA` as **Connect As**. Then, click **Login**.

The first time you log in as a new user, the licensing page appears.

- g) To acknowledge this information, click **I agree** at the bottom of the page.

The Database Home page appears.

- 3. Using Enterprise Manager, view information about the `EXAMPLE` tablespace. Answer the following questions about it:
 - a) In Enterprise Manager, select **Administration** > Database Administration > Storage > **Tablespaces**.
 - b) Click the `EXAMPLE` tablespace name.
 - c) *Question 1:* What percentage of free space can be used up before the Warning threshold is reached?

Answer: 85%

Tablespace Full Metric Thresholds

Space Used (%)

This tablespace is using the database default space used thresholds.

Warning (%) **85**

Critical (%) **97**

Solutions for Practice 5: Managing Database Storage Structures (continued)

- d) From the **Actions** drop-down list, select **Show Tablespace Contents**, and then click **Go**.
- e) The Show Tablespace Contents page appears.

Previous 1-10 of 418 Next 10			
Segment Name	Type	Size (KB)	Extents
SH.CUSTOMERS	TABLE	12,288	27
SH.SUPPLEMENTARY_DEMOGRAPHICS	TABLE	4,096	19
OE.PRODUCT_DESCRIPTIONS	TABLE	2,072	18

- f) *Question 2:* How many segments are there in the EXAMPLE tablespace?

Answer: 418

- g) Select **INDEX** from the **Type** drop-down list in the Search region, and then click **Go**.

Previous 1-10 of 77 Next 10			
Segment Name	Type	Size (KB)	Extents
SH.CUSTOMERS_PK	INDEX	1,024	16
OE.PROD_NAME_IX	INDEX	512	8
OE.PROD_DESC_PK	INDEX	320	5

- h) *Question 3:* Which index in the EXAMPLE tablespace takes up the most space?

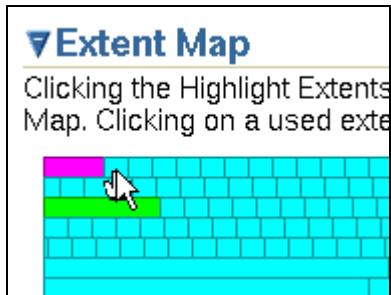
Answer: SH.CUSTOMERS_PK

Question 4: Which segment is stored physically first in the tablespace? That is, which one is stored right after the tablespace header?

- i) Scroll to the bottom of the page, and then click the plus icon to the left of the Extent Map label.
- j) After several seconds, the extent map appears. Note that the map legend indicates that pink is the tablespace header.
- k) Scroll back to the top of the page, select **All Types** from the **Type** drop-down list, and then click **Go**.

Solutions for Practice 5: Managing Database Storage Structures (continued)

- l) Click the extent just to the right of the tablespace header extent.



- m) Scroll to the top of the page again, and note the segment that is being pointed to:

<input type="radio"/>	PM.SYS_LOB0000051820C00004\$\$\$	LOBSEGMENT	128	<u>2</u>
<input type="radio"/>	PM.SYS_LOB0000051819C00015\$\$\$	LOBSEGMENT	128	<u>2</u>
<input checked="" type="radio"/>	HR.REGIONS	TABLE	64	<u>1</u>
<input type="radio"/>	HR.LOCATIONS	TABLE	64	<u>1</u>

Answer: HR.REGIONS

4. Create a new, locally managed tablespace (LMT) called INVENTORY of size 5 MB.

Note: The INVENTORY tablespace will be used in later practice sessions during this class.

- In Enterprise Manager, select **Administration** > Database Administration > Storage > **Tablespaces**.
- Click **Create**.
- Enter INVENTORY as the tablespace name, and verify that **Extent Management** is **Locally Managed**, **Type** is **Permanent**, **Status** is **Read Write**, and **Use Bigfile tablespace** is not selected.
- Click **Add** in the Datafiles region.
- On the Add Datafile page, enter inventory01.dbf for **File Name**, and 5 MB as **File Size**.
- Click **Continue**.
- Click the **Storage** tab, and verify that **Extent Allocation** is **Automatic**, **Segment Space Management** is **Automatic**, and **Logging** is enabled.
- Click the **General** tab.
- Click **Show SQL** to see the SQL that will be run, and then click **Return**.
- Click **OK**, and a successful Update Message appears.

Solutions for Practice 5: Managing Database Storage Structures (continued)

5. As the DBA1 user, run the lab_05_05.sql script to create and populate a table (called X) in the INVENTORY tablespace. What error do you eventually see?

- a) Right-click the desktop and select **Open Terminal** to start a command shell. Then enter:

```
$ cd labs
```

- b) Enter the following to run the script:

```
$ sqlplus dba1/oracle @lab_05_05.sql
```

- c) Note that there is eventually an error ORA-01653 stating that the table cannot be extended. There is not enough space to accommodate all of the rows to be inserted.

```
1024 rows created.

SQL> insert into x select * from x
  2 /
insert into x select * from x
*
ERROR at line 1:
ORA-01653: unable to extend table DBA1.X by 128 in tablespace INVENTORY

SQL> commit
  2 /

Commit complete.
```

6. Go to the Enterprise Manager window and define space for 50 MB in the tablespace instead of 5 MB, while keeping the same single data file in the tablespace. What is the ALTER statement that is executed to make this change?
- a) Select **Administration** > Database Administration > Storage > **Tablespaces**.
- b) Select the INVENTORY tablespace, and then click **Edit**.
- c) In the **Datafiles** region, click **Edit**.
- d) Change **File Size** from 5 MB to 50 MB.
- e) Click **Continue** to return to the General tabbed page.
- f) Click **Show SQL** to see the SQL that will be run. Note that it is an ALTER DATABASE statement. Click **Return**.

Solutions for Practice 5: Managing Database Storage Structures (continued)

Show SQL

Return

```
ALTER DATABASE DATAFILE  
'/u01/app/oracle/oradata/ORCL/datafile/inventory01.dbf'  
RESIZE 50M
```

- g) Click **Apply**.
7. Go back to the terminal window and run the `lab_05_07.sql` script. It drops the table and reexecutes the original script that previously returned the space error.
- a) Go to the SQL*Plus window.
- b) Enter the following to run the script:

```
$ sqlplus dba1/oracle @lab_05_07.sql
```

- c) Note that the same number of row inserts are attempted, and there is no error because of the increased size of the tablespace.
8. In an terminal window, run the `lab_05_08.sql` script to clean up the tablespace for later practice sessions.
- a) Enter the following to run the script:

```
$ sqlplus dba1/oracle @lab_05_08.sql
```

Solutions for Practice 6: Administering User Security

Background: You need to create a user account for Jenny Goodman, the new human resources department manager. There are also two new clerks in the human resources department, David Hamby and Rachel Pandya. All three of them must be able to log in to the ORCL database, select data from, and update records in the HR.EMPLOYEES table. The manager also needs to be able to insert and delete new employee records. Ensure that if the new users forget to log out at the end of the day, they will automatically be logged out after 15 minutes. You also need to create a new user account for the inventory application that you are installing.

1. **Mandatory task:** Review and run the lab_06_01.sh script (located in the /home/oracle/labs directory) to create the INVENTORY user, which you will use in the next practice.

a) In a terminal window, enter:

```
cd /home/oracle/labs
more lab_06_01.sh
./lab_06_01.sh
```

2. Create a profile named HRPFILE, allowing 15 minutes idle time.
 - a) Invoke Enterprise Manager as the DBA1 user in the SYSDBA role for your ORCL database.
 - b) Select **Administration** > Schema > Users & Privileges > **Profiles**.
 - c) Click the **Create** button.
 - d) Enter **HRPROFILE** in the **Name** field.
 - e) Enter **15** in the **Idle Time (Minutes)** field.

Solutions for Practice 6: Administering User Security (continued)

Create Profile

General Password

* Name

Details

CPU/Session (Sec./100)	<input type="text" value="DEFAULT"/>	
CPU/Call (Sec./100)	<input type="text" value="DEFAULT"/>	
Connect Time (Minutes)	<input type="text" value="DEFAULT"/>	
Idle Time (Minutes)	<input type="text" value="15"/>	

Database Services

Concurrent Sessions (Per User)	<input type="text" value="DEFAULT"/>	
Reads/Session (Blocks)	<input type="text" value="DEFAULT"/>	
Reads/Call (Blocks)	<input type="text" value="DEFAULT"/>	
Private SGA (KBytes)	<input type="text" value="DEFAULT"/>	
Composite Limit (Service Units)	<input type="text" value="DEFAULT"/>	

Show SQL Cancel OK

- f) Leave all the other fields set to **DEFAULT**.
 - g) Click the **Password** tab, and review the Password options, which are currently all set to **DEFAULT**.
 - h) Optionally, click the **Show SQL** button, review your underlying SQL statement, and then click **Return**.
 - i) Finally, click **OK** to create your profile.
3. Set the **RESOURCE_LIMIT** initialization parameter to **TRUE** so that your profile limits will be enforced.
- a) Select **Administration** > Database Administration > Database Configuration > **All Initialization Parameters**.
 - b) Enter **RESOURCE_LIMIT** in the **Filter** field, and then click **Go**.
 - c) Select **TRUE** from the **Value** drop-down list, and then click **Apply**.

Solutions for Practice 6: Administering User Security (continued)

Initialization Parameters

Current SPFile

The parameter values listed here are currently used by the running instance(s). You can change static parameters in SPFile mode.

Name: Basic: Modified: Dynamic: Category:

Filter on a name or partial name

☐ Apply changes in current running instance(s) mode to SPFile. For static parameters, you must restart the database.

Name	Help	Revisions	Value	Comments	Type	Basic	Modified	Dynamic	Category
resource_limit			TRUE		Boolean			<input checked="" type="checkbox"/>	Resource Manager

4. Create the role named HRCLERK with SELECT and UPDATE permissions on the HR.EMPLOYEES table.
 - a) Select **Administration** > Schema > Users & Privileges > **Roles**.
 - b) Click the **Create** button at the top-right of the page.
 - c) Enter **HRCLERK** in the **Name** field. This role is not authenticated.
 - d) Click **Object Privileges**.
 - e) Select **Table** from the **Select Object Type** drop-down list, and then click **Add**.
 - f) Enter **HR.EMPLOYEES** in the **Select Table Objects** field.
 - g) Move the SELECT and UPDATE privileges to the **Selected Privileges** box. Click **OK**.
 - h) Click the **Show SQL** button, and review your underlying SQL statement.

```
CREATE ROLE "HRCLERK" NOT IDENTIFIED
GRANT SELECT ON "HR"."EMPLOYEES" TO "HRCLERK"
GRANT UPDATE ON "HR"."EMPLOYEES" TO "HRCLERK"
```

- i) Click **Return**, and then click **OK** to create the role.

Solutions for Practice 6: Administering User Security (continued)

Create Role

Show SQL Cancel OK

General Roles System Privileges **Object Privileges** Consumer Groups Switching Privileges

Select Object Type Function Add

Delete

Select	Object Privilege	Schema	Object
<input checked="" type="checkbox"/>	SELECT	HR	EMPLOYEES
<input checked="" type="checkbox"/>	UPDATE	HR	EMPLOYEES

5. Create the role named HRMANAGER with INSERT and DELETE permissions on the HR.EMPLOYEES table. Grant the HRCLERK role to the HRMANAGER role.
 - a) Select **Administration** > Schema > Users & Privileges > **Roles**.
 - b) Click **Create**.
 - c) Enter **HRMANAGER** in the **Name** field. This role is not authenticated.
 - d) Click **Object Privileges**.
 - e) Select **Table** from the **Select Object Type** drop-down list, and then click **Add**.
 - f) Enter **HR.EMPLOYEES** in the **Select Table Objects** field.
 - g) Move the INSERT and DELETE privileges to the **Selected Privileges** box. Click **OK**.
 - h) Click **Roles**, and then click **Edit List**.
 - i) Move the HRCLERK role into the **Selected Roles** box, and then click **OK**.

Create Role

Show SQL Cancel OK

General **Roles** System Privileges Object Privileges Consumer Groups Switching Privileges

Edit List

Role	Admin Option
HRCLERK	<input type="checkbox"/>

- j) Click the **Show SQL** button, and review your underlying SQL statement.

```
CREATE ROLE "HRMANAGER" NOT IDENTIFIED
GRANT DELETE ON "HR"."EMPLOYEES" TO "HRMANAGER"
GRANT INSERT ON "HR"."EMPLOYEES" TO "HRMANAGER"
GRANT "HRCLERK" TO "HRMANAGER"
```

- k) Click **Return**, and then click **OK** to create the role.

Solutions for Practice 6: Administering User Security (continued)

6. Create an account for David Hamby, a new HR clerk.
 - a) Select **Administration** > Schema > Users & Privileges > **Users**.
 - b) Click **Create**, and enter **DHAMBY** in the Name field.
 - c) Select **HRPROFILE** for the Profile.
 - d) Select Password Authentication, and enter **newuser** as password. Enter it into the **Confirm Password** field also. Select the **Expire Password now** check box so that David will have to change the password the first time he logs in.

Create User

Show SQL Cancel OK

General Roles System Privileges Object Privileges Quotas Consumer Groups Switching Privileges Proxy Users

* Name

Profile

Authentication

* Enter Password

* Confirm Password

For Password choice, the role is authorized via password.

☒ Expire Password now

Default Tablespace

Temporary Tablespace

Status ☐ Locked ☒ Unlocked

- e) Click **Roles**. Notice that the CONNECT role has automatically been assigned to the user.
 - f) Add the HRCLERK role by clicking **Edit List** and moving the HRCLERK role into the **Selected Roles** box. Click **OK** to close the Modify Roles window.

Edit User: DHAMBY

Show SQL Revert Apply

General **Roles** System Privileges Object Privileges Quotas Consumer Groups Proxy Users

Modify

Role	Admin Option	Default
CONNECT	<input type="checkbox"/>	<input checked="" type="checkbox"/>
HRCLERK	<input type="checkbox"/>	<input checked="" type="checkbox"/>

- g) Click **OK** again to create the user.
7. Create an account for Rachel Pandya, another new HR clerk. Repeat the steps under step 6 with RPANDYA as the username.

Solutions for Practice 6: Administering User Security (continued)

8. Create an account for Jenny Goodman, the new HR manager. Repeat the steps under step 6 with JGOODMAN as the username and selecting the HRMANAGER role instead of the HRCLERK role.

a) Click the **Show SQL** button and review your underlying SQL statement.

```
CREATE USER "JGOODMAN" PROFILE "HRPROFILE" IDENTIFIED BY "*****" PASSWORD
EXPIRE ACCOUNT UNLOCK
GRANT "CONNECT" TO "JGOODMAN"
GRANT "HRMANAGER" TO "JGOODMAN"
```

b) Click **Return**, and then click **OK** to create the user.

9. Test the new users in SQL*Plus. Connect to the ORCL database as the DHAMBY user. Use oracle as the new password. Select the row with EMPLOYEE_ID=197 from the HR.EMPLOYEES table. Then, attempt to delete it. (You should get the “insufficient privileges” error.)

a) In a terminal window, enter:

```
sqlplus dhamby/newuser
```

or if you are already in SQL*Plus, use the CONNECT command. If you reconnect as dhamby in SQL*Plus, the login and change-of-password session looks like this:

```
SQL> CONNECT dhamby/newuser

ERROR:
ORA-28001: the password has expired

Changing password for dhamby
New password: oracle    <<<Password does not appear on screen
Retype new password: oracle    <<<Password does not appear on screen
Password changed

Connected to:
Oracle Database 10g Enterprise Edition Release 10.2.0.1.0 - Production
With the Partitioning, OLAP and Data Mining options

SQL>
```

b) Select the salary for EMPLOYEE_ID=197 from the HR.EMPLOYEES table.

```
SQL> SELECT salary FROM hr.employees WHERE EMPLOYEE_ID=197;

      SALARY
-----
        3000
```

Solutions for Practice 6: Administering User Security (continued)

c) Now attempt to delete the same record from the `hr.employees` table.

```
SQL> DELETE FROM hr.employees WHERE EMPLOYEE_ID=197;
DELETE FROM hr.employees WHERE EMPLOYEE_ID=197
      *
ERROR at line 1:
ORA-01031: insufficient privileges
```

10. Repeat the delete attempt as the JGOODMAN user. After deleting the row, issue a rollback, so that you still have the original 107 rows.

a) Connect to the ORCL database as the JGOODMAN user.

```
SQL> connect jgoodman/newuser
ERROR:
ORA-28001: the password has expired
<Change the password as shown above>
```

b) Select the row with `EMPLOYEE_ID=197` from the `HR.EMPLOYEES` table.

```
SQL> SELECT salary FROM hr.employees WHERE EMPLOYEE_ID=197;

      SALARY
-----
      3000
```

c) Now delete the same row from the `HR.EMPLOYEES` table.

```
SQL> DELETE FROM hr.employees WHERE EMPLOYEE_ID=197;

1 row deleted.
```

d) Roll back the delete operation (because this was just a test).

```
SQL> rollback;

Rollback complete.
```

e) Confirm that you still have 107 rows in this table.

```
SQL> SELECT COUNT(*) FROM hr.employees;

      COUNT(*)
-----
           107

SQL>
```

Question 1: Where was the row stored after deletion?

Answer: It was stored in the Undo tablespace.

Solutions for Practice 6: Administering User Security (continued)

Question 2: When you created the new users, you did not select a default or temporary tablespace. What determines the tablespaces that the new users will use?

Answer: The system-defined default permanent and temporary tablespaces

Question 3: You did not grant the CREATE SESSION system privilege to any of the new users, but they can all connect to the database. Why?

Answer: Because Enterprise Manager automatically assigns the CONNECT role to the new users, and CREATE SESSION is contained within that role

11. Review the `lab_05_01.sql` script and the `lab_05_01.txt` log file that it generated when you created the DBA1 user.

- a) Double-click the **oracle's Home** icon on your desktop.
- b) Navigate to the **labs** directory.
- c) Double-click the **lab_05_01.sql** file, and review its content.
- d) When you have finished reviewing the file, click the **Up** icon to return to the **labs** directory.
- e) Now, double-click the **lab_05_01.txt** file, which was created when you executed the `lab_05_01.sql` file.
- f) When you have finished reviewing the file, click the X (Close Window) icon.

12. Use SQL*Plus to connect to the ORCL database as the RPANDYA user. Change the password to `oracle`. (You must change the password, because this is the first connection as RPANDYA.) Leave RPANDYA connected during the next lesson or at the end of the day. HRPROFILE specifies that users whose sessions are inactive for more than 15 minutes will automatically be logged out. Verify that the user was automatically logged out by trying to select from the `HR.EMPLOYEES` table again.

```
ERROR at line 1:
ORA-02396: exceeded maximum idle time, please connect again
```

Solutions for Practice 7: Managing Schema Objects

Background: You need to create schema objects for the new inventory application. Work as the DBA1 user in the SYSDBA role for your ORCL database.

1. Return to the Enterprise Manager browser session, or invoke EM as the DBA1 user in the SYSDBA role for your ORCL database.
2. In the INVENTORY tablespace, create the PRODUCT_MASTER table in the INVENTORY schema. The specifications of the table are:

PRODUCT_ID number(7). This is the primary key field. (Constraint name: PK_INV)
PRODUCT_NAME varchar2(50) with a Not NULL constraint
CODE varchar2(10) with a Not NULL constraint
REORDER_THRESHOLD number(5) with a check constraint ensuring that the number is always greater than zero (Constraint name: CHK_REORDER)
COST number(5,2)
PRICE number(5,2)

- a) Select **Administration** > Schema > Database Objects > **Tables**.
- b) Click **Create**.
- c) Leave the default table type to be **Standard, Heap Organized permanent**, and then click **Continue**.
- d) Enter PRODUCT_MASTER as table name and set **Schema** to INVENTORY, and leave the Tablespace setting as the default for the user, which is INVENTORY.
- e) Enter the first five columns' information (including Data Type, Size, Scale, and Not NULL); but no constraints yet.
- f) Click **Add 5 Table Columns**, and then enter the last column, PRICE, along with its type and size.
- g) Click the **Constraints** tab.
- h) Set the drop-down list value to PRIMARY, and then click **Add**.
- i) Give the constraint a name by entering PK_INV in the **Name** field.
- j) Double-click PRODUCT_ID in the left list to move it to the right list, to make it alone the primary key. Then, click **Continue**.

Solutions for Practice 7: Managing Schema Objects (continued)

- k) Set the constraint type drop-down list value to CHECK, and then click **Add**.
- l) Enter CHK_REORDER for the name of the check constraint.
- m) Enter “reorder_threshold > 0” (without the quotation marks) in the **Check Condition** field.

Definition

Name

CHK_REORDER

* Check Condition

reorder_threshold > 0

Attributes

☐ Disabled

☐ Deferrable - In subsequent transactions this allows constraint checking to be deferred until the end of the transaction.

☐ Initially Deferred - Set the default deferred behavior to check constraints at the end of a transaction.

☒ Validate - Check to ensure all existing data meets the constraint criteria.

☐ Do not enforce the constraint (RELY) - Constraint is not used to enforce data integrity. It is used to express the relationship between tables and views.

- n) Leave the Attributes settings at their default settings, and then click **Continue**.
- o) Your constraint list should now look like this:

Select	Name	Type	Table Columns	Disabled	Deferrable	Initially Deferred	Validate	RELY
	PK_INV	PRIMARY	PRODUCT_ID	NO	NO	NO	YES	NO
	CHK_REORDER	CHECK		NO	NO	NO	YES	NO

- p) Click **OK** to create the table. If you receive errors, correct them, and then click **OK** again.

Solutions for Practice 7: Managing Schema Objects (continued)

3. In the INVENTORY tablespace, create the PRODUCT_ON_HAND table in the INVENTORY schema. You have been given the lab_07_03.sql script to create the table, but there is a problem with it (*intentionally created to enhance your learning experience*). Fix the problem, and run the script. If you cannot find the error right away, then go ahead and run the original script in SQL*Plus to see the error message. This helps you discover and solve the problem. The specifications of the table are:

PRODUCT_ID number(7). This field should have a foreign key constraint linking it to the PRODUCT_ID field in the PRODUCT_MASTER table.

QUANTITY number(5)

WAREHOUSE_CITY varchar2(30)

LAST_UPDATE date

- a) Edit the lab_07_03.sql in the /home/oracle/labs directory. It has an error in it. If you can spot the error, make the change to correct it. Run the script to create the table by entering this on the OS command line:

```
$ sqlplus dba1/oracle @lab_07_03.sql
```

- b) The error in the script is that “(PRODUCT_ID)” is missing after “FOREIGN KEY.” So, add “(PRODUCT_ID)”.

```
SQL> CREATE TABLE INVENTORY.PRODUCT_ON_HAND
 2  (
 3  PRODUCT_ID NUMBER(7),
 4  QUANTITY NUMBER(5),
 5  WAREHOUSE_CITY VARCHAR2(30),
 6  LAST_UPDATE DATE,
 7  CONSTRAINT FK_PROD_ON_HAND_PROD_ID
 8  FOREIGN KEY REFERENCES
 9  INVENTORY.PRODUCT_MASTER (PRODUCT_ID) VALIDATE
10 )
11 /
      FOREIGN KEY REFERENCES
          *
```

ERROR at line 8:
ORA-00906: missing left parenthesis

- c) Remove the prompt commands:

```
prompt There is an error in this statement. It will not
prompt run successfully unless fixed.
```

- d) Run the script. The table should be created without error.

Solutions for Practice 7: Managing Schema Objects (continued)

4. In the INVENTORY tablespace, create the OBSOLETE_PRODUCTS table in the INVENTORY schema. This table definition is very much like that of the PRODUCT_MASTER table, so you can use Enterprise Manager's ability to "Define Using SQL" rather than using "Column Specification." The specifications of the table are:

PRODUCT_ID number(7). This is the primary key field.
PRODUCT_NAME varchar2(50) with a Not Null constraint
CODE varchar2(20) with a Not Null constraint
COST number(5,2)
PRICE number(5,2)

- In Enterprise Manager, select **Administration** > Schema > Database Objects > **Tables**.
- Click **Create**.
- Leave the default setting for **Table Organization**, and then click **Continue**.
- Enter OBSOLETE_PRODUCTS for **Name**.
- Enter INVENTORY for **Schema**.
- Leave the Tablespace setting as the default for this schema.
- Set the **Define Using** drop-down list to **SQL**.
- In the SQL region, enter the following statement:

```
SELECT product_id, product_name, code, cost, price
FROM inventory.product_master
```

Create Table

Show SQL Cancel OK

General Constraints Storage Options

* Name

Schema

Tablespace

Organization **Standard, Heap Organized**

Define Using

SQL

Enter a SQL select statement below. The results of this query will be used to populate this table with data.

CREATE TABLE AS SELECT product_id, product_name, code, cost, price FROM inventory.product_master

- Click **OK** to create the table.

Solutions for Practice 7: Managing Schema Objects (continued)

5. In the INVENTORY tablespace, create an index called OBS_CODE on the CODE column of the OBSOLETE_PRODUCTS table in the INVENTORY schema. Choose an appropriate index type: either B-tree or Bitmap. Explain the reason for your choice.

a) *Question:* Which type of index is appropriate, and why?

Answer: B-tree, because the CODE column can contain many different values, not just a small, finite set.

- b) Select **Administration** > Schema > Database Objects > **Indexes**, and then click **Create**.
- c) Enter OBS_CODE as **Name**.
- d) Enter INVENTORY as **Schema**.
- e) Enter INVENTORY.OBSOLETE_PRODUCTS as **Table Name**.
- f) Click **Populate Columns**. The column names have been filled into the table. If your browser prompts you to remember entered values, click **No**.
- g) Enter an ORDER value of 1 for the CODE column.

Table Columns			
Column Name	Data Type	Sorting Order	Order
PRODUCT_ID	NUMBER	ASC ▾	
PRODUCT_NAME	VARCHAR2	ASC ▾	
CODE	VARCHAR2	ASC ▾	1
COST	NUMBER	ASC ▾	
PRICE	NUMBER	ASC ▾	
Add Column Expression			

- h) Click **Show SQL**, confirm the SQL statement looks like the statement below, and then click **Return**.

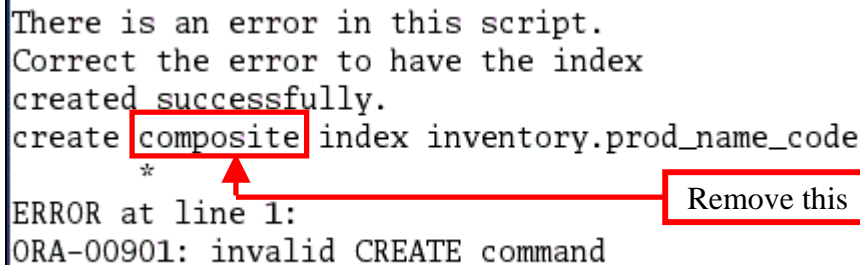
```
CREATE INDEX "INVENTORY"."OBS_CODE" ON
"INVENTORY"."OBSOLETE_PRODUCTS" ("CODE")
```

- i) Click **OK** to create the index.

Solutions for Practice 7: Managing Schema Objects (continued)

6. In the INVENTORY tablespace, create an index called PROD_NAME_CODE on the combined PRODUCT_NAME and CODE columns of the PRODUCT_MASTER table. Use the lab_07_06.sql script (*which contains an error, intentionally created to enhance your learning experience*). Correct the error and run the script. If you cannot find the error right away, then run the original script in SQL*Plus to see the error message. This will help you discover and solve the problem.
- a) Edit the lab_07_06.sql in the /home/oracle/labs directory. It has an error in it. If you can spot the error, make the change to correct it.

```
There is an error in this script.
Correct the error to have the index
created successfully.
create composite index inventory.prod_name_code
      *
ERROR at line 1:
ORA-00901: invalid CREATE command
```



The corrected statement is:

```
create index inventory.prod_name_code
on inventory.product_master(product_name,code)
```

- b) Remove the prompt commands:

```
prompt  There is an error in this script.
prompt  Correct the error to have the index
prompt  created successfully.
```

- c) Run the script to create the index by entering the following on the OS command line:

```
$ sqlplus dba1/oracle @lab_07_06.sql
```

- d) The error in the script is that it should not have the word “COMPOSITE”. Delete the word, run the script, and the index should be created without error.

7. In the INVENTORY tablespace, use SQL*Plus to create a combined index on the PRODUCT_ID and QUANTITY columns of the PRODUCT_ON_HAND table. The index name should be POH_PROD_ID_QTY.

- a) At the SQL*Plus prompt, enter the following command:

```
SQL> create index inventory.poh_prod_id_qty on
inventory.product_on_hand(product_id, quantity);
```

Solutions for Practice 7: Managing Schema Objects (continued)

8. You receive an update for the inventory application that requires you to add two columns to the PRODUCT_MASTER table. Add a column called PRIMARY_SOURCE of the data type varchar2(50). Add another column called SECONDARY_SOURCE of the data type varchar2(50). What is the SQL that executes to do this?

- a) Select **Administration** > Schema > Database Objects > **Tables**.
- b) Enter INVENTORY in the **Schema** field, and then click **Go**.
- c) Select the PRODUCT_MASTER table, and then click **Edit**.
- d) Enter PRIMARY_SOURCE in the **Name** field under PRICE, and set **Data Type** to VARCHAR2, and **Size** to 50.
- e) Enter SECONDARY_SOURCE in the next available **Name** field, and set **Data Type** to VARCHAR2, and **Size** to 50.
- f) Click **Show SQL**.

```
ALTER TABLE "INVENTORY"."PRODUCT_MASTER"  
ADD ( "PRIMARY_SOURCE" VARCHAR2(50),  
"SECONDARY_SOURCE" VARCHAR2(50))
```

- g) Click **Return**, and then click **Apply**.
9. You receive another update for the inventory application. This change request instructs you to drop the OBSOLETE_PRODUCTS table and add the OBSOLETE column to the PRODUCT_MASTER table, with data type DATE. Do this using EM. What clause is added to the end of the DROP TABLE statement to also remove the table constraints?

- a) Select **Administration** > Schema > Database Objects > **Tables**.
- b) Enter INVENTORY in the **Schema** field, and then click **Go**.
- c) Select the OBSOLETE_PRODUCTS table, and then click **Delete With Options**.
- d) Keep the default settings for performing DROP, and then click **Show SQL** to see the CASCADE CONSTRAINTS option.

```
DROP TABLE "INVENTORY"."OBSOLETE_PRODUCTS" CASCADE CONSTRAINTS
```

- e) Click **Return**, and then click **Yes**.
- f) Select the PRODUCT_MASTER table, and then click **Edit**.

Solutions for Practice 7: Managing Schema Objects (continued)

g) In the first empty **Name** field, enter OBSOLETE, and set **Data Type** to DATE.

h) Click **Apply**.

10. Another change request to the inventory application instructs you to create a view called WAREHOUSE_VW (Navigation aid: Administration > Views). The view is in the INVENTORY schema and displays (in this order):

- The name of the product
- The amount of “product on hand”
- The warehouse city name

You will have to join two tables together to create this view.

a) Select **Administration** > Schema > Database Objects > **Views**.

b) Click **Create**.

c) Enter WAREHOUSE_VW in the **Name** field, and INVENTORY in the **Schema** field.

d) Enter the following in the **Query Text** field:

```
select product_name, quantity, warehouse_city
from product_master pm, product_on_hand poh
where pm.product_id = poh.product_id
```

Create View

General Options Object

* Name WAREHOUSE_VW

* Schema INVENTORY

Aliases

☐ Replace the view if exists

* Query Text select product_name, quantity, warehouse_city
from product_master pm, product_on_hand poh
where pm.product_id = poh.product_id

e) Click **OK**.

Solutions for Practice 7: Managing Schema Objects (continued)

11. You receive a notice from developers that there is a certain type of query that will be run very frequently, and they want to be sure that it runs in less than one second. You need to run the query and see how it performs. First, run the `lab_07_11_a.sql` script to generate some test data of the volume indicated in the request. Then, run the query in the `lab_07_11_b.sql` script several times to see the average run time. Note that it takes several seconds to run each time. Create a function-based index on the `CODE` column that will improve the performance of this query.

- a) Right-click the desktop and select **Open Terminal**. Then enter the following at the command prompt to populate the table with test data. This will take three to five minutes to run.

```
$ cd /home/oracle/labs
$ sqlplus dba1/oracle @lab_07_11_a.sql
```

- b) Enter the following on the command line to run the test query.

```
SQL> @lab_07_11_b.sql
```

- c) Enter `/` (a slash), and press [Enter] to run it again. Repeat this several times until you establish an average run time. It should take several seconds each time.
- d) Select **Administration** > Schema > Database Objects > **Indexes**.
- e) Click **Create**.
- f) Enter `CODE_FUNC` in the **Name** field.
- g) Enter `INVENTORY` in the **Schema** field.
- h) Enter `INVENTORY.PRODUCT_MASTER` in the **Table Name** field.
- i) Click **Populate Columns**.
- j) Click **Add Column Expression**.
- k) Enter `upper(substr(code,5,2))` in the newly added empty **Column Name** field.
- l) Enter `1` in the **Order** field beside the expression that you just entered.
- m) Click **Show SQL** to confirm that the SQL statement looks like this:

```
CREATE INDEX "INVENTORY"."CODE_FUNC" ON "INVENTORY"."PRODUCT_MASTER"
(upper(substr(code,5,2)))
```

Solutions for Practice 7: Managing Schema Objects (continued)

- n) Click **Return**.
- o) Click **OK**.
- p) Return to the SQL*Plus command line and run the benchmark query a few more times. Note that the execution time is greatly reduced.

12. Use *iSQL*Plus* to identify the data dictionary view name that you would use to list all constraints that the INVENTORY user can see.

- a) In the Enterprise Manager browser session, select **Database** > Related Links > **iSQL*Plus**.

The *iSQL*Plus* Connection Role page appears.

- b) Select **Normal**, and then click **Continue**.
- c) On the Login page, enter inventory as **Username** and verysecure as the **Password**, leave the **Connect Identifier** set to its default value, and then click **Login**.

The Confirm window appears, with the Password Manager offering to remember this login for you. Click **No**.

- d) You are looking for things that the INVENTORY user has access to. So, you must use a view with the ALL_ prefix. You can assume that the name of the view that shows constraint information begins with "ALL_CON." In the *iSQL*Plus* Workspace, enter:

```
SELECT * FROM dictionary
WHERE table_name like 'ALL_CON%'
ORDER BY table_name;
```

- e) Click the **Execute** button.

Solutions for Practice 7: Managing Schema Objects (continued)

The screenshot shows the iSQL*Plus Workspace interface. At the top right, it says "Connected as HR@orcl". The main area has a text input field with the prompt "Enter SQL, PL/SQL and SQL*Plus statements." and a "Clear" button. Below the input field, the following SQL query is entered:

```
SELECT * FROM dictionary
WHERE table_name like 'ALL_CON%'
ORDER BY table_name;
```

Below the query, there are four buttons: "Execute", "Load Script", "Save Script", and "Cancel". Below the buttons, a table displays the results of the query:

TABLE_NAME	COMMENTS
ALL_CONSTRAINTS	Constraint definitions on accessible tables
ALL_CONS_COLUMNS	Information about accessible columns in constraint definitions
ALL_CONS_OBJ_COLUMNS	List of types an object column or attribute is constrained to in the tables accessible to the user
ALL_CONTEXT	Description of all active context namespaces under the current session

- f) Review the returned rows. One of them has the COMMENTS value “Constraints definitions on accessible tables.” The object name is ALL_CONSTRAINTS.
13. How many indexes are owned by the INVENTORY user? You are looking for database objects owned by the INVENTORY user, so you know that the data dictionary view begins with the “USER_” prefix. The view name is USER_INDEXES.

- a) In the iSQL*Plus Workspace, enter:

```
SELECT * FROM user_indexes;
```

- b) Click the **Execute** button.

How many indexes are owned by the INVENTORY user?

Answer: 4

Solutions for Practice 7: Managing Schema Objects (continued)

Connected as INVENTORY@orcl

Workspace

Enter SQL, PL/SQL and SQL*Plus statements.

SELECT * FROM user_indexes;

Clear

Execute Load Script Save Script Cancel

INDEX_NAME	INDEX_TYPE	TABLE_OWNER	TABLE_NAME	TABL
CODE_FUNC	FUNCTION-BASED NORMAL	INVENTORY	PRODUCT_MASTER	TABL
PROD_NAME_CODE	NORMAL	INVENTORY	PRODUCT_MASTER	TABL
PK_INV	NORMAL	INVENTORY	PRODUCT_MASTER	TABL
POH_PROD_ID_QTY	NORMAL	INVENTORY	PRODUCT_ON_HAND	TABL

- c) Note that there are four rows returned.
- d) Log out of *iSQL*Plus* by clicking **Logout** in the top-right corner of the window. Then, close the window by clicking **X** (Close icon) in the top-right corner.

Solutions for Practice 8: Managing Data and Concurrency

Background: The Help desk just received a call from Susan Mavris, an HR representative, complaining that the database is “frozen.” Upon questioning the user, you find that she was trying to update John Chen’s personnel record with his new phone number, but when she entered the new data, her session froze and she could not do anything else. SQL script files are provided for you in the /home/oracle/labs directory.

1. Make an uncommitted update to the row in question by running the lab_08_01.sql script. Do not worry if the session seems to “hang”—this is the condition you are trying to create.
 - a) Enter the following to run the script. When the script completes executing, you will see a note stating that an uncommitted update has been made.

```
$ sqlplus dba1/oracle @lab_08_01.sql
```

```
SQL> show user
USER is "NGREENBERG"
SQL> update hr.employees set phone_number='650.555.1212' where employee_id = 110;

1 row updated.

SQL> prompt User "ngreenberg" made an update and left it uncommitted in this session.
User "ngreenberg" made an update and left it uncommitted in this session.
SQL>
SQL>
SQL>
```

2. Make an attempt to update the same row in a separate session by running, in a separate terminal window, the lab_08_02.sql script. Make sure you see the message “Update is being attempted now” before moving on.
 - a) Right-click the desktop and select **Open Terminal** to start another command shell. Then enter the following to run the second script.

```
$ sqlplus dba1/oracle @lab_08_02.sql
```

```
Sleeping for 20 seconds to ensure first process gets the lock first.

PL/SQL procedure successfully completed.

Sleep is finished.
Connected.
USER is "SMAVRIS"
Update is being attempted now.
```

Solutions for Practice 8: Managing Data and Concurrency (continued)

3. Using the **Blocking Sessions** link on the Performance page, detect which session is causing the locking conflict.
 - a) In Enterprise Manager, click the Performance page.
 - b) The first time you invoke the Performance page, a Software License Agreement is displayed. Press [A] to accept and continue.
 - c) Click **Blocking Sessions** in the **Additional Monitoring Links** area.

Blocking Sessions											
Page Refreshed May 31, 2005 5:07:44 PM											
Expand All Collapse All View Session Kill Session											
Select	Username	Sessions Blocked	Session ID	Session Serial Number	SQL Hash Value	Wait Class	Wait Event	P1	P2	P3	Seconds in Wait
	Blocking Sessions										
	NGREENBERG	1	133	1784		Idle	SQL*Net message from client	1650815232	1	0	538
	SMAVRIS	0	158	11895	6smgtv6h8958b	Application	eng: TX - row lock contention	1415053318	327682	540	538

4. What was the last SQL statement that the blocking session executed?
 - a) Select the NGREENBERG session, and then click **View Session**.
 - b) Click the hash value link named “**Previous SQL**.”

Application	
Current SQL	None
Current SQL Command	UNKNOWN
Previous SQL	6smgtv6h8958b
Last Call Elapsed Time	21 Minutes, 13 Seconds
SQL Trace	DISABLED
Open Cursors	0
Program	sqlplus@EDRSR9P1 (TNS V1-V3)
Service	SYS\$USERS
Current Module	SQL*Plus
Current Action	Unavailable

Solutions for Practice 8: Managing Data and Concurrency (continued)

- c) Note the SQL that was most recently run.

Text

```
update hr.employees set phone_number='650.555.1212'  
where employee_id = 110
```

5. Resolve the conflict in favor of the user who complained, by killing the *blocking* session. What SQL statement resolves the conflict?
- a) Click the browser's **Back** button.
 - b) Now, on the Session Details: NGREENBERG page, click **Kill Session**.
 - c) Leave the Options set to **Kill Immediate**, and then click **Show SQL** to see the statement that is going to be executed to kill the session.

Note: Your session and serial number are most likely to be different from those shown here.

```
ALTER SYSTEM KILL SESSION '133,1784' IMMEDIATE
```

- d) Click **Return**, and then click **Yes** to carry out the KILL SESSION command.
6. Return to the SQL*Plus command window, and note that SMAVRIS's update has now completed successfully. It may take a few seconds for the success message to appear.

```
USER is "SMAVRIS"  
Update is being attempted now.  
  
1 row updated.  
  
Update is completed.  
SQL> █
```

- a) Close all open SQL sessions by entering `exit`, and then close the terminal windows.

Solutions for Practice 9: Managing Undo Data

Background: A new version of your application will include several reports based on very long-running queries. Configure your system to support these reports.

1. Use the Undo Advisor to calculate the amount of undo space required to support a report that takes two days to run, on the basis of an analysis period of the last seven days.
 - a) In Enterprise Manager, select **Administration** > Related Links > **Advisor Central**.
 - b) Click **Undo Management**.
 - c) Click **Undo Advisor**.
 - d) In the **New Undo Retention** field, enter 2. Then select **days** from the drop-down list.
 - e) Select Last Seven Days in the **Analysis Time Period** drop-down list. Results of the analysis will be displayed.

Note: The values that you see are likely to be different from those shown here.

Advisor	
* New Undo Retention	2 days
Analysis Time Period	Last Seven Days
Choose the time period that best represents system activity	
Selected Analysis Time Period	October 21, 2005 8:00 AM To October 28, 2005 8:00 AM
Analysis	
Required Tablespace Size for New Undo Retention (MB)	110
Required Tablespace Size for Current Undo Retention (MB)	11
Maximum Extensible Undo Tablespace Size (MB)	32652
Undo Retention (days)	15
Best Possible Undo Retention (days)	594

- f) *Question:* What does the analysis recommend as “Required Tablespace Size for New Undo Retention”?

Answer: 110 MB

- g) Click **OK**.

Solutions for Practice 9: Managing Undo Data (continued)

2. Resize the undo tablespace to support the retention period required by the new reports (or 1 GB, whichever is smaller). Do this by increasing the size of the existing data file.

a) Click **Edit Undo Tablespace**.

b) *Question:* What are the two ways to add space to a tablespace?

Answer: Add a new data file or increase the size of an existing data file.

c) To increase the size of the existing data file, select the data file, and then click **Edit**.

d) Enter the new value from step 1 (f) in the **File Size** field. Round off to the nearest 100 MB (not to exceed 1 GB). Make sure the box for “Automatically extend datafile when full” is not selected. Then, click **Continue**.

e) Click **Show SQL** to confirm that the SQL statement looks similar to this:

```
ALTER DATABASE DATAFILE
'/u01/app/oracle/oradata/ORCL/datafile/o1_mf_undotbs1_1pzfonlk_.dbf'
RESIZE 200M
ALTER DATABASE DATAFILE
'/u01/app/oracle/oradata/ORCL/datafile/o1_mf_undotbs1_1pzfonlk_.dbf'
AUTOEXTEND OFF
```

f) Click **Return**, then click **Apply** to implement your change.


Solutions for Practice 10: Implementing Oracle Database Security

Background: You have just been informed of suspicious activities in the HR . JOBS table in your ORCL database. All maximum salaries seem to fluctuate in a strange way. You decide to enable standard database auditing and monitor data manipulation language (DML) activities in this table.

Log in as the DBA1 user (with oracle password, connect as SYSDBA) and perform the necessary tasks either through Enterprise Manager Database Control or through SQL*Plus. All scripts for this practice are in the /home/oracle/labs directory.

1. Use Enterprise Manager to enable database auditing. Set the AUDIT_TRAIL parameter to XML.
 - a) Invoke Enterprise Manager as the DBA1 user in the SYSDBA role for your ORCL database.
 - b) Select **Administration** > Schema > Users & Privileges > **Audit Settings**.

Audit Settings

 Standard auditing for this database is disabled. To enable it, change the Audit Trail configuration parameter and restart the database. All the audit trail information either have no data or contain only obsolete audit records.

Configuration	Audit Trails
Audit Trail NONE	Database Audit Trail Audited Failed Logins
Audit SYS User Operations FALSE	Audited Privileges
Audit File Directory /u01/app/oracle/admin/orcl/adump	Audited Objects

Audit File Directory value is effective only when Audit Trail is set to "OS" or "XML".

- c) Click the **NONE** link.
- d) On the Initialization Parameters page, click the **SPFile** tab.
- e) Enter audit in the **Name** field and then click **Go**.
- f) For the **audit_trail** parameter, select the **XML** value.
- g) Click **Show SQL**.

Show SQL

```
ALTER SYSTEM SET audit_trail = "XML" SCOPE=SPFILE
```

- h) Review the statement and then click **Return**.
- i) On the Initialization Parameters page, click **Apply**.

Solutions for Practice 10: Implementing Oracle Database Security (continued)

- Because you changed a static parameter, you must restart the database. Do so by running the `lab_10_02.sh` script.

- In a terminal window, enter:

```
cd /home/oracle/labs
./lab_10_02.sh
```

Continue with the next step when you see that the database is restarted and the script has exited out of SQL*Plus.

- Back in Enterprise Manager, select `HR.JOBS` as the audited object and `DELETE`, `INSERT`, and `UPDATE` as **Selected Statements**. Gather audit information by session.

- Click the Database home page tab to ensure that Enterprise Manager had time to update the status of the database and its agent connections. Because the database has been restarted, you have to log in to EM again as the `DBA1` user.
- Then, select **Administration** > Schema > Users & Privileges > **Audit Settings**.

Audit Settings

Audit information can be located in the database or in an OS file. Some information is always written to the OS audit file. Other information can optionally be written to either the OS audit file or to the database.

Configuration

Audit Trail [XML](#)
Audit SYS User Operations [FALSE](#)
Audit File Directory [/u01/app/oracle/admin/orcl/adump](#)
Audit File Directory value is effective only when Audit Trail is set to "OS" or "XML".

Default Options For Future Audited Objects [0](#)

Audit Trails

Database Audit Trail [Audited Failed Logins](#)
[Audited Privileges](#)
[Audited Objects](#)

Audited Privileges (0) **Audited Objects (0)** **Audited Statements (0)**

Privilege User Proxy Search

Select	Privilege	User	Proxy	Success	Failure
No object found.					

Add

- Click the **Audited Objects** tab, and then click the **Add** button.
- On the Add Audited Object page, ensure that the **Object Type** is Table, and enter `HR.JOBS` in the **Table** field (or use the flashlight icon to retrieve this table).
- Move `DELETE`, `INSERT`, and `UPDATE` into the **Selected Statements** area by double-clicking each of them.
- Click **Show SQL**.

Show SQL

AUDIT DELETE, INSERT, UPDATE ON HR.JOBS BY SESSION

Solutions for Practice 10: Implementing Oracle Database Security (continued)

- g) Review the statement, and then click **Return**.
 - h) Click **OK** to activate this audit.
4. Provide input for the audit, by executing the `lab_10_04.sh` script. This script creates the `AUDIT_USER` user, connects to `SQL*Plus` as this user, and multiplies the values in the `MAX_SALARY` column by 10. Then, the `HR` user connects and divides the column values by 10. Finally, the `AUDIT_USER` user is dropped again.

- a) In a terminal window, enter:

```
cd /home/oracle/labs
./lab_10_04.sh
```


5. In Enterprise Manager, review the audited objects.
- a) Select **Administration** > Schema > Users & Privileges > **Audit Settings**.
 - b) Click **Audited Objects** in the **Audit Trails** area, which is on the right side of the page.
 - c) On the Audited Objects page, review the collected information, and optionally click **Show SQL**.

Audited Objects				
<div>Filter Result Return</div> <div>▼ Hide SQL</div> <div>SELECT "OBJECT_SCHEMA", "OBJECT_NAME", "DB_USER", "STATEMENT_TYPE", "EXTENDED_TIMESTAMP" FROM SYS.DBA_COMMON_AUDIT_TRAIL WHERE (action between 1 and 16) or (action between 19 and 29) or (action between 32 and 41) or (action = 43) or (action between 51 and 99) or (action = 103) or (action between 110 and 113) or (action between 116 and 121) or (action between 123 and 128) or (action between 160 and 162)</div>				
Schema	Object Name	User Name	Action	Time (In Session's Time Zone)
HR	JOBS	AUDIT_USER	SESSION REC	2005-10-21 17:52:33.783793000 -7:0
HR	JOBS	HR	SESSION REC	2005-10-21 17:52:34.147582000 -7:0

- d) Click **Return**.
6. Undo your audit settings for `HR.JOBS`, disable database auditing, and then restart the database by using the `lab_10_06.sh` script.
- a) On the Audit Settings page, click the **Audited Objects** tab.
 - b) Enter `HR` as **Schema**, and then click **Search**.
 - c) Select all three rows, and then click **Remove**.

Solutions for Practice 10: Implementing Oracle Database Security (continued)


- d) On the Confirmation page, click **Show SQL**.

 **Confirmation**

NoYes

Are you sure you want to remove the 3 selected audited objects?

The audited statements you remove will no longer be audited on the objects.

 [Hide SQL](#)

NOAUDIT DELETE ON HR.JOBS
NOAUDIT INSERT ON HR.JOBS
NOAUDIT UPDATE ON HR.JOBS

- e) Review the statements, then click **Yes** to confirm your removal.
- f) On the Audit Settings page, click **XML** in the **Configuration** region.
- g) On the Initialization Parameters page, click the **SPFile** tab.
- h) On the SPFile page, enter `audit` in the **Name** field, and then click **Go**.
- i) For the `audit_trail` parameter, select the **NONE** value.
- j) Click **Show SQL**.

Show SQL

ALTER SYSTEM SET audit_trail = "NONE" **SCOPE=SPFILE**

- k) Review the statement, and then click **Return**.
- l) On the Initialization Parameters page, click **Apply**.
- m) Because you changed a static parameter, you must restart the database. Do so by running the `lab_10_06.sh` script. In a terminal window, enter:

```
cd /home/oracle/labs
./lab_10_06.sh
```

7. Maintain your audit trail: Because you are completely finished with this task, delete all audit files from the `/u01/app/oracle/admin/orcl/adump` directory.

- a) In a terminal window, enter:

```
cd /u01/app/oracle/admin/orcl/adump
ls
rm -f *
```

- b) Close the terminal window.

Solutions for Practice 11: Configuring the Oracle Network Environment

Background: Users need to connect to your ORCL database. Work with them to enable connections by using different methods. Ensure that users can use connect-time failover to take advantage of a backup listener.

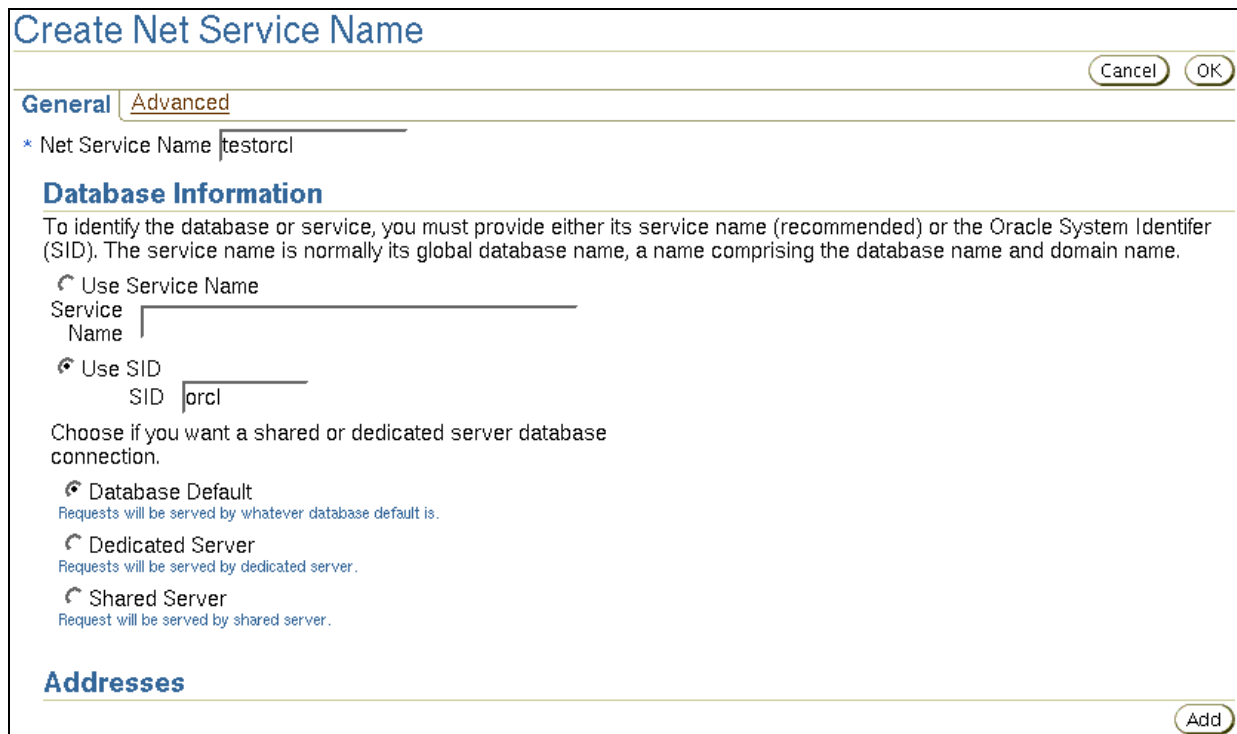
1. Make a copy of your `listener.ora` and `tnsnames.ora` files. They are in the `$ORACLE_HOME/network/admin` directory.
 - a) In an terminal window, enter `cd $ORACLE_HOME/network/admin` to navigate to the `/u01/app/oracle/product/10.2.0/db_1/network/admin` directory.
 - b) Enter `cp listener.ora listener.old` to create a copy of the `listener.ora` file.
 - c) Enter `cp tnsnames.ora tnsnames.old` to create a copy of the `tnsnames.ora` file.
 - d) Enter `ls -l`, if you want to see the copies and their privileges in your directory.
2. Navigate to the Net Services Administration page. Start by clicking the **Listener** link on the Database home page.
 - a) Invoke Enterprise Manager as the DBA1 user in the SYSDBA role for your ORCL database.
 - b) On the Database Instance – Home page, click the **Listener** link in the **General** region.
 - c) In the **Related Links** region, click **Net Services Administration**.
3. Modify your local Names Resolution file so that you can connect to another database. Name the connection to a partner's ORCL database **testorcl**.
 - a) On the Net Services Administration page, select **Local Naming** from the Administer drop-down list, and then click **Go**.

The **Netservices Administration: Host Login** page appears.

- b) If you previously saved the `oracle` username and `oracle` password as preferred credentials for your host login, then they appear on the screen. If not, enter **oracle** as Username and Password, select the **Save as Preferred Credential** check box, and then click **Login**.
- c) On the Local Naming page, click **Create** to enter a new network service name.
- d) Enter **testorcl** as Net Service Name.
- e) Select **Use SID**, and enter **orcl** as SID.

Solutions for Practice 11: Configuring the Oracle Network Environment (continued)

f) Select **Database Default**.



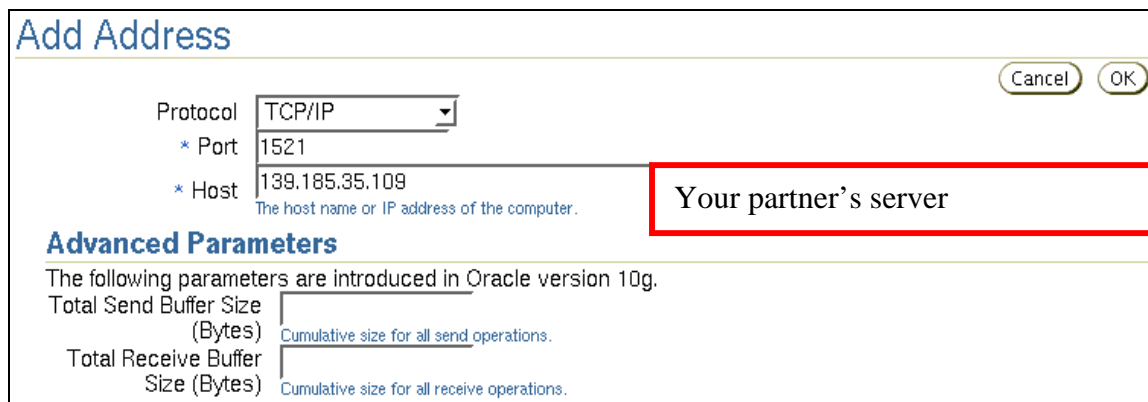
The 'Create Net Service Name' dialog box has two tabs: 'General' and 'Advanced'. The 'General' tab is active. It contains a text field for 'Net Service Name' with the value 'testorcl'. Below this is the 'Database Information' section, which explains that the service name is normally its global database name. It offers two radio button options: 'Use Service Name' (with a text field for 'Service Name') and 'Use SID' (with a text field for 'SID' containing 'orcl'). The 'Use SID' option is selected. Below these are three radio button options for the connection type: 'Database Default' (selected), 'Dedicated Server', and 'Shared Server'. Each option has a brief description. At the bottom is the 'Addresses' section with an 'Add' button.

g) Click **Add** in the **Addresses** region.

Click **No** if you are asked, “Do you want to remember the values you filled in?”

h) On the Add Address page, select the following values:

- Protocol: **TCP/IP**
- Port: **1521**
- Host: <Enter the host name or IP address of **your partner's computer**. It could be something like `edrsr9p1.us.oracle.com` or like `139.185.35.109`>



The 'Add Address' dialog box has 'Cancel' and 'OK' buttons at the top right. It contains three fields: 'Protocol' (a dropdown menu set to 'TCP/IP'), 'Port' (a text field with '1521'), and 'Host' (a text field with '139.185.35.109'). The 'Host' field is highlighted with a red rectangle, and the text 'Your partner's server' is written next to it. Below these fields is the 'Advanced Parameters' section, which states that the following parameters are introduced in Oracle version 10g. It includes two text fields: 'Total Send Buffer Size (Bytes)' and 'Total Receive Buffer Size (Bytes)', both with descriptions of their cumulative sizes.

Solutions for Practice 11: Configuring the Oracle Network Environment (continued)

- i) Click **OK** to return to the Create Net Service Name properties page.
- j) Click **OK**.

The Creation Message appears: Net Service “testorcl” created successfully.

- 4. In Enterprise Manager, test access to your partner’s ORCL database as the **system** user with the **oracle** password by using the **testorcl** Local Naming.
 - a) Select **testorcl** on the Local Naming page, and then click **Test Connection**.

The “Test Connection To Net Service Name: testorcl” appears.

- b) Enter **system** as Username and **oracle** as Password, and then click **Test**.

The Processing page displays status information. Then, it is followed by a success message. *If you receive any errors or warnings, resolve them.*

- 5. Test your changes to the network configuration by using SQL*Plus or iSQL*Plus. Again, use: `system/oracle@testorcl`. To see your partner’s information, select the `instance_name` and `host_name` columns from the `v$instance` table.

- a) In an terminal window, enter:

```
sqlplus system/oracle@testorcl
```

The Oracle SQL*Plus window opens. *If you receive any errors or warnings, resolve them.*

- b) At the SQL> prompt, enter the following command:

```
select instance_name, host_name from v$instance;
```

Solutions for Practice 11: Configuring the Oracle Network Environment (continued)

```
[oracle@EDRSR30P1 oracle]$ sqlplus system/oracle@testorcl

SQL*Plus: Release 10.2.0.0.0 - Beta on Mon May 16 16:38:46 2005

Copyright (c) 1982, 2004, Oracle. All rights reserved.

Connected to:
Oracle Database 10g Enterprise Edition Release 10.2.0.0.0 - Beta
With the Partitioning, OLAP and Data Mining options

SQL> select instance_name, host_name from v$instance;

INSTANCE_NAME
-----
HOST_NAME
-----
orcl
EDRSR9P1

SQL> █
```

You should see your partner's host name.

- c) Enter **exit** to exit your SQL*Plus session.
6. Create a LISTENER2 listener to support connect-time failover. Use port 1561 for this listener. First, log out of Enterprise Manager and run the `lab_11_06.sh` script to configure the `NetProperties` file.
 - a) Logout of Enterprise Manager.
 - b) Run the `lab_11_06.sh` script at the operating system command prompt.
 - c) Log in to Enterprise Manager as the DBA1 user in the SYSDBA role for your ORCL database.
 - d) On the Database Instance – Home page, click the **Listener** link in the **General** region.
 - e) In the **Related Links** region, click **Net Services Administration**.
 - f) On the Net Services Administration page, select **Listeners** from the **Administer** drop-down list, and then click **Go**. Enter host credentials as `oracle` and `oracle` for username and password, and then click **Login**.

Solutions for Practice 11: Configuring the Oracle Network Environment (continued)

- g) On the Listeners page, which gives you an overview of the existing listeners, click the **Create** button.

The Create Listener page appears.

- h) Enter **LISTENER2** as Listener Name, and then click **Add** to add a listener address.
- i) Enter or confirm the following values:
- Protocol: TCP/IP
 - Port: **1561**
 - Host: <Your computer's host name; for example, `edrsr30p1.us.oracle.com`>
- j) Click **OK**.

Create Listener

Cancel OK

General Authentication Logging & Tracing Static Database Registration Other Services

* Listener Name LISTENER2

Addresses

Listener must have at least one address. If address is changed, listener will be stopped before applying changes.

Add Edit Remove

Select	Protocol	Protocol Details
<input checked="" type="radio"/>	TCP/IP	Host edrsr30p1.us.oracle.com Port 1561

- k) Click the **Static Database Registration** tab, and then click the **Add** button to connect the new listener with your ORCL database.
- l) Enter the following values:
- Service Name: **orcl**
 - Oracle Home Directory: **/u01/app/oracle/product/10.2.0/db_1**
 - Oracle System Identifier (SID): **orcl**

Add Database Service

Cancel OK


* Service Name orcl

* Oracle Home Directory /u01/app/oracle/product/10.2.0/db_1

* Oracle System Identifier (SID) orcl

- m) Click **OK** to add the database service.
- n) Click **OK** to create the LISTENER2 listener.

**Solutions for Practice 11: Configuring the Oracle Network Environment
(continued)**


 **Creation Message**

Listener "LISTENER2" created successfully.

Listeners: /u01/app/oracle/product/10.2.0/db_1/network/admin

Listener Name Go

Edit Delete Actions Start/Stop Go Create

Select	Name	Protocol Details	Status	Enterprise Manager Target
	<u>LISTENER2</u>	Protocol TCP/IP Host edrsr30p1.us.oracle.com Port 1561	Stopped	Not a target

7. Start the LISTENER2 listener.
- a) Confirm that the **LISTENER2** listener and **Start/Stop** Actions are selected, and then click **Go**.
 - b) Click **OK** on the Start/Stop page.

A confirmation message appears with a **View Details** link.

- c) Optionally, click the **View Details** link, review the listener status information, and use the **Back** icon of your browser to return to the previous page.

Solutions for Practice 12: Proactive Maintenance

Background: You want to proactively monitor your ORCL database so that common problems can be fixed before they affect users. This practice session invents some issues so that you can familiarize yourself with the tools that are available. First, execute scripts to set up your Automatic Database Diagnostic Management (ADDM) environment.

1. Create a new, locally managed tablespace called TBSADDM. Its addm1.dbf data file is 50 MB. Ensure that the TBSADDM tablespace does not use Automatic Segment Space Management (ASSM). Execute the lab_12_01.sh script to perform these tasks.

a) In a terminal window, enter:

```
cd /home/oracle/labs
./lab_12_01.sh
```

2. Create a new ADDM user, identified by ADDM. Assign the TBSADDM tablespace as default tablespace. Assign the TEMP tablespace as temporary tablespace. Grant the following roles to the ADDM user: CONNECT, RESOURCE, and DBA. Execute the lab_12_02.sh script to perform these tasks.

a) In a terminal window, enter:

```
./lab_12_02.sh
```

3. Use the DBMS_ADVISOR package to set the database activity time to 30 minutes. As an ADDM user, drop and create the ADDM table and gather statistics for this table. Create a snapshot in Automatic Workload Repository (AWR). Execute the lab_12_03.sh script to perform these tasks.

a) In a terminal window, enter:

```
./lab_12_03.sh
```


4. Create activity to be analyzed. Execute the lab_12_04.sh script to perform these tasks.

a) In a terminal window, enter the following. You may have to press [Enter] after you see that eight PL/SQL procedures have completed, in order to see the command prompt again.

```
./lab_12_04.sh
```

Solutions for Practice 12: Proactive Maintenance (continued)

5. In Enterprise Manager, review the Performance page as a user connected as SYSDBA. View performance data in real time with a 15-seconds refresh cycle. After a while, you should see a spike on the “Average Active Sessions” graph. This is your activity to be analyzed. Looking at the graph, you can already determine that this instance is suffering from concurrency problems.




Note: Depending on when you run the workload, you may see differences between your graph and the one provided as a possible solution.

After the spike is finished, execute the `lab_12_05.sh` script. This script forces the creation of a new snapshot and gathers statistics on your ADDM table.

- Invoke Enterprise Manager as the DBA1 user in the SYSDBA role for your ORCL database.
- Click the **Performance** tab.

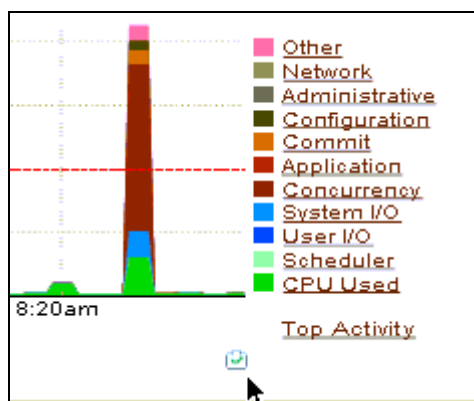
If this is the first time that you accessed the Performance page, you need to accept the Adobe license agreement. Follow the directions in the pop-up window to accept the agreement.

Solutions for Practice 12: Proactive Maintenance (continued)



c) After the spike has finished, in a terminal window, enter:

```
./lab_12_05.sh
```



Note: The icon next to the cursor in this screenshot is a shortcut to the Automatic Database Diagnostic Monitor (ADDM) page.

Solutions for Practice 12: Proactive Maintenance (continued)

6. Look at the **Performance Analysis** findings in order of their impact. There are several access paths to this information.

Looking at the Performance Analysis section, you see that the first finding (in the SQL Tuning Recommendations category) has a 100% impact on the system. So your first impulse is to look at this finding in more detail. However, looking at this SQL statement does not yet help you to understand the concurrency problem of your database.

Research the next finding under Schema Recommendations: **Read and write contention of database blocks was consuming significant database time**. Here you are advised to use the Automatic Segment Space Management (ASSM) feature for your ADDM table.

- a) Navigate to the Database home page, and then click **Advisor Central** at the bottom of the page.

Advisor Central

Page Refreshed Oct 25, 2005 8:46:32 AM PDT [Refresh](#)

Advisors

[ADDM](#) [Memory Advisor](#) [MTTR Advisor](#)
[Segment Advisor](#) [SQL Access Advisor](#) [SQL Tuning Advisor](#)
[Undo Management](#)

Advisor Tasks [Change Default Parameters](#)

Search

Select an advisory type and optionally enter a task name to filter the data that is displayed in your results set.

Advisory Type Task Name Advisor Runs Status

All Types Last Run All [Go](#)

By default, the search returns all uppercase matches beginning with the string you entered. To run an exact or case-sensitive match, double quote the search string. You can use the wildcard symbol (%) in a double quoted string.

Results

[View Result](#) [Delete](#) [Actions](#) [Re-schedule](#) [Go](#)

Select	Advisory Type	Name	Description	User	Status	Start Time	Duration (seconds)	Expires In (days)
<input checked="" type="radio"/>	ADDM	ADDM:1099806762 1 161	ADDM auto run: snapshots [160, 161], instance 1, database id 1099806762	ADDM	COMPLETED	Oct 25, 2005 8:26:51 AM	2	30
<input type="radio"/>	Segment Advisor	SYS_AUTO_SPCADV 70525102005	Auto Space Advisor	SYS	COMPLETED	Oct 24, 2005 10:00:09 PM	5	29

- b) Your ADDM task should already be displayed. If not, search for it and display it on this page.
- c) Select the task, and then click the **View Result** button (or alternatively, click the name of the task).

Solutions for Practice 12: Proactive Maintenance (continued)

Performance Analysis		
Task Name	ADDM:1099806762_1_161	Time Range Oct 25, 2005 8:11:00 AM to Oct 25, 2005 8:41:00 AM
Database Time (minutes)	3.2	Period Start Time Oct 25, 2005 8:00:22 AM PDT
Task Owner	ADDM	Average Active Sessions 0.1
		Period Duration (minutes) 26.4
		View Snapshots View Report
Impact (%)	Finding	Recommendations
100	SQL statements consuming significant database time were found.	3 SQL Tuning
25.6	Read and write contention on database blocks was consuming significant database time.	3 Schema
22.4	A hot data block with concurrent read and write activity was found. The block belongs to segment "ADDM.ADDM" and is block 45 in file 6.	1 Application Analysis 1 Schema
13	Wait event "latch: In memory undo latch" in wait class "Concurrency" was consuming significant database time.	1 Application Analysis
12.6	Time spent on the CPU by the instance was responsible for a substantial part of database time.	2 SQL Tuning
5.1	Hard parsing of SQL statements was consuming significant database time.	
3.9	Waits on event "log file sync" while performing COMMIT and ROLLBACK operations were consuming significant database time.	1 Application Analysis 1 Host Configuration
2.4	Database latches in the "Other" wait class were consuming significant database time.	
2.2	Contention on the high watermark (HW) enqueue was consuming significant database time.	
► Informational Findings		


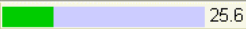
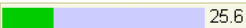
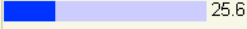
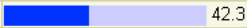
Looking at the Performance Analysis section, you see that the first finding has a 100% impact on the system. So your first impulse is to look at this finding in more detail.

d) Click the **SQL statements consuming significant database time were found** link.

Performance Finding Details		
Database Time (minutes)	3.2	Period Start Time Oct 25, 2005 8:00:22 AM PDT
Task Owner	ADDM	Task Name ADDM:1099806762_1_161
		Period Duration (minutes) 26.4
		Average Active Sessions 0.1
Finding	SQL statements consuming significant database time were found.	
Impact (minutes)	3.2	
Impact (%)	100	
Recommendations		
Schedule SQL Tuning Advisor		
Select All Select None Show All Details Hide All Details		
Select Details	Category	Benefit (%)
<input type="checkbox"/> Hide	SQL Tuning	76.5
Action Tune the PL/SQL block with SQL_ID "8brpf00mu7ubh" . Refer to the "Tuning PL/SQL Applications" chapter of Oracle's "PL/SQL User's Guide and Reference" SQL Text declare t number, begin for t in 1..2222 loop insert into addm values (Null,'a')... SQL ID 8brpf00mu7ubh		
Rationale SQL statement with SQL_ID "8brpf00mu7ubh" was executed 8 times and had an average elapsed time of 18 seconds.		
<input type="checkbox"/> Show	SQL Tuning	58
<input checked="" type="checkbox"/> Show	SQL Tuning	3.5
Findings Path		
Expand All Collapse All		
Findings	Impact (%)	Additional Information
SQL statements consuming significant database time were found.	100	

Solutions for Practice 12: Proactive Maintenance (continued)

- e) Review the recommendations on the Performance Finding Details page. However, looking at this SQL statement does not yet help you to understand the concurrency problem of your database. Click the **Back** icon in your Web browser.
- f) On the Automatic Database Diagnostic Monitor (ADDM) page, click the **Read and write contention of database blocks was consuming significant database time** link . This finding appears as type **Schema** under the **Recommendations** heading.

Performance Finding Details			
Database Time (minutes)	3.2	Period Start Time	Oct 25, 2005 8:00:22 AM PDT
Task Owner	ADDM	Task Name	ADDM:1099806762_1_161
		Period Duration (minutes)	26.4
		Average Active Sessions	0.1
Finding	Read and write contention on database blocks was consuming significant database time.		
Impact (minutes)	0.8		
Impact (%)		25.6	
Recommendations			
Show All Details Hide All Details			
Details	Category	Benefit (%) ▾	
▼ Hide	Schema		25.6
Action	Consider using ORACLE's recommended solution of automatic segment space management in a locally managed tablespace for the tablespace "TBSADDM" containing the TABLE "ADDM.ADDM" with object id 52818. Alternatively, you can move this object to a different tablespace that is locally managed with automatic segment space management.		
Database Object	ADDM.ADDM		
Rationale	There was significant read and write contention on TABLE "ADDM.ADDM" with object id 52818.		
Database Object	ADDM.ADDM		
▶ Show Schema		25.6	
▶ Show Schema		25.6	
Findings Path			
Expand All Collapse All			
Findings	Impact (%)	Additional Information	
▼ Read and write contention on database blocks was consuming significant database time.		25.6	
Wait class "Concurrency" was consuming significant database time.		42.3	

- g) You are advised to use the Automatic Segment Space Management feature for your ADDM table.
7. To implement the recommendation, you must re-create the object. Create a new, locally managed tablespace, called TBSADDM2 with a 50 MB datafile, called addm2_1.dbf. Ensure that the TBSADDM2 tablespace uses the the Automatic Segment Space Management feature. Then, execute the lab_12_07.sh script to drop the ADDM table, to re-create it in the new tablespace, to gather statistics and to take a new snapshot.
 - a) In Enterprise Manager, select **Administration** > Database Administration > Storage > **Tablespaces**.
 - b) Click **Create**.
 - c) Enter TBSADDM2 as the tablespace name, and verify that **Extent Management** is **Locally Managed**, **Type** is **Permanent**, **Status** is **Read Write**, and **Use Bigfile tablespace** is not selected.

Solutions for Practice 12: Proactive Maintenance (continued)

- d) Click **Add** in the **Datafiles** region.
- e) On the Add Datafile page, enter addm2_1.dbf for **File Name**, and 50 MB as **File Size**.
- f) Click **Continue**.
- g) Click the **Storage** tab, and verify that **Extent Allocation** is **Automatic**, **Segment Space Management** is **Automatic**, and **Logging** is enabled.
- h) Click the **General** tab.
- i) Click **Show SQL** to see the SQL that will be run, and then click **Return**.
- j) Click **OK**, and a successful Update Message appears.
- k) In a terminal window, enter:

```
./lab_12_07.sh
```

- 8. Execute your workload again. (The lab_12_08.sh script is identical to the lab_12_04.sh script.)
 - a) In a terminal window, enter the following. You may have to press [Enter] after you see that eight PL/SQL procedures have completed, in order to see the command prompt again.

```
./lab_12_08.sh
```


- 9. In Enterprise Manager, review the Performance page as a user connected as SYSDBA. View performance data in real time with a 15-seconds refresh cycle. After a while, you should see a spike on the “Average Active Sessions” graph.

Note: Depending on when you run the workload, you may see differences between your graph and the one provided as a possible solution.

After the spike is finished, execute the lab_12_09.sh script. (The lab_12_09.sh script is identical to the lab_12_05.sh script.) This script forces the creation of a new snapshot and gathers statistics on your ADDM table.

- a) Invoke Enterprise Manager as the DBA1 user in the SYSDBA role for your ORCL database.
- b) Click on the **Performance** tabbed page.

Solutions for Practice 12: Proactive Maintenance (continued)



c) After the spike is finished, enter the following in a terminal window:

```
./lab_12_09.sh
```

10. Review the Performance Analysis on the Database home page.

- Navigate to the Database home page, and then click **Advisor Central** at the bottom of the page.
- Click the topmost ADDM task name.

Performance Analysis			
Task Name	ADDM:1101003407_1_67		Time Range
Database Time (minutes) 3.2		Period Start Time Nov 4, 2005 4:46:48 AM PST	Period Duration (minutes) 3.4
Task Owner	ADDM	Average Active Sessions	0.9
		View Snapshots View Report	
Impact (%)	Finding	Recommendations	
100	SQL statements consuming significant database time were found.	3 SQL Tuning	
16.3	Time spent on the CPU by the instance was responsible for a substantial part of database time.	3 SQL Tuning	
11.6	Wait event "latch: redo copy" in wait class "Configuration" was consuming significant database time.	1 Application Analysis	
7.1	Wait event "latch: In memory undo latch" in wait class "Concurrency" was consuming significant database time.	1 Application Analysis	
6.4	Wait event "latch: enqueue hash chains" in wait class "Other" was consuming significant database time.	1 Application Analysis	

You see that there are no longer any schema-related recommendations. By moving the ADDM table to the locally managed TBSADDM2 tablespace, which uses the Automatic Autoextend Segment feature, you obviously fixed the root cause of this problem.

Solutions for Practice 12: Proactive Maintenance (continued)

11. To not affect other practice session, execute the `lab_12_11.sh` script to clean up your environment.

a) In a terminal window, enter:

```
./lab_12_11.sh
```

Solutions for Practice 13: Performance Management

Background: Users are complaining about slower-than-normal performance for operations involving the human resources and order-entry applications. When you question other members of the DBA staff, you find that maintenance was recently performed on some of the tables belonging to the HR schema. You need to troubleshoot and make changes as appropriate to resolve the performance problems. SQL script files are provided for you in the `/home/oracle/labs` directory. Other directories are individually named.

1. Log in to SQL*Plus as the DBA1 user and perform maintenance on tables in the HR schema by running the `lab_13_01.sql` script.

- a) Connect to the database as the DBA1 user and perform some maintenance on the `employees` table by running the `lab_13_01.sql` script.

```
$ cd labs
$ sqlplus dba1/oracle
SQL> @lab_13_01.sql
```

2. You get calls from HR application users saying that a particular query is taking longer than normal to execute. The query is in the `lab_13_02.sql` script. As the HR user, run it.

- a) Enter the following in SQL*Plus , while logged in as the HR user:

```
SQL> @lab_13_02.sql
```

3. Using Enterprise Manager, locate the HR session in which the above statement was just executed, and view the execution plan for that statement.

- a) In Enterprise Manager, click **Search Sessions** on the Performance tabbed page.

Solutions for Practice 13: Performance Management (continued)

- b) Change the Filter for the search criteria to “DB User,” enter HR in the field to the right of that, and then click **Go**.

Database Instance: [orcl.oracle.com](#) > Search Sessions Logged in As DBA1

Search Sessions

Search

☒ Specify search criteria

Filter DB User Go

For DB User, the search returns all uppercase matches. To run an exact or case-sensitive match, double quote the search string. For other filters, the search returns all case-sensitive matches. You can always use the wildcard symbol (%).

☐ Specify search criteria using WHERE clause [Display Columns](#)

Go

(Example: SID > 5 AND USERNAME LIKE 'SCOTT')

Results

SID	DB User	Program	Service	Module	Action	Client	Machine	OS User
159	HR	sqlplus@EDRSR9P1 (TNS V1-V3)	SYS\$USERS	SQL*Plus			EDRSR9P1	oracle

- c) Click the **SID** number in the **Results** listing.
- d) You now see the Session Details page for this session. Click the hash value link to the right of the **Previous SQL** label.

Database Instance: [orcl.oracle.com](#) > [Top Activity](#) > Session Details: HR (159) Logged in As DBA1

Session Details: HR (159)

Collected From Target **Oct 28, 2005 1:49:13 PM** View Data | Real Time: 15 Second Refresh Refresh

Kill Session Enable SQL Trace Disable SQL Trace

General [Activity](#) [Statistics](#) [Open Cursors](#) [Blocking Tree](#) [Wait Event History](#)

Server	Client	Application
Current Status INACTIVE Serial Number 2114 DB User Name HR OS Process ID 14577 Logged On Since Oct 28, 2005 1:36:20 PM Logged On For 12 Minutes, 53 Seconds Connection Type DEDICATED Type USER Resource Consumer Group Unavailable	OS User Name oracle OS Process ID 14576 Host EDRSR9P1 Terminal pts/4 Current Client ID Unavailable Current Client Info Unavailable	Current SQL None Current SQL Command UNKNOWN Previous SQL bckcqw5pd108f Last Call Elapsed Time 12 Minutes, 45 Seconds SQL Trace DISABLED Open Cursors 42 Program sqlplus@EDRSR9P1 (TNS V1-V3) Service SYS\$USERS Current Module SQL*Plus Current Action Unavailable

Solutions for Practice 13: Performance Management (continued)

- e) That shows you the SQL Details page for the last SQL statement executed by that session, which is the one in question. Click the **Plan** tab to see the execution plan for the query.

Database Instance: [orcl.oracle.com](#) > [Top Activity](#) > SQL Details: bckcqw5pd108f Logged in As DBA1

SQL Details: bckcqw5pd108f

Switch to SQL ID View Data Real Time: Manual Refresh

Text

```
select * from hr.employees where employee_id = 200
```

Details

Select the plan hash value to see the details below. Plan Hash Value

[Statistics](#) [Activity](#) **[Plan](#)** [Tuning Information](#)

Data Source **Cursor Cache** Capture Time **Oct 28, 2005 1:52:40 PM** Parsing Schema **HR** Optimizer Mode **ALL_ROWS**

[Expand All](#) | [Collapse All](#)

Operation	Object	Object Type	Order	Rows	Size (KB)	Cost	Time (sec)	CPU Cost	I/O Cost
SELECT STATEMENT			2			3			
TABLE ACCESS FULL	EMPLOYEES	TABLE	1	1	0.066	3	1	57207	3

- f) You see in the Operation column that this query is doing a full table scan (TABLE ACCESS FULL). Because you know that the query's condition is an equality comparison on the primary key (EMPLOYEE_ID), you decide to investigate the status of the primary key index.
4. Using Enterprise Manager, check to see the status of the EMPLOYEE table's index on EMPLOYEE_ID. See if it is VALID.
- Select **Administration** > Schema > Database Objects > **Indexes**.
 - Select Table Name as the **Search By** value.
 - Enter HR as **Schema Name**.
 - Enter employees for **Object Name**.
 - Click **Go**, and the list of six indexes appears.
 - Click the index named EMP_EMP_ID_PK.
 - Click the index link and note that the status of the index is UNUSABLE.

Solutions for Practice 13: Performance Management (continued)

5. Now that you have seen one index with a non-VALID status, you decide to check all indexes. Using SQL*Plus, as the HR user, find out which HR schema indexes do not have STATUS of VALID. To do this, you can query a data dictionary view with a condition on the STATUS column.

- a) Go to the SQL*Plus session where you are still logged in as the HR user, and run this query:

```
SQL> select index_name, table_name, status
       from user_indexes where status <> 'VALID';
```

- b) See that the output lists six indexes, all on the EMPLOYEES table.

```
SQL> select index_name, table_name, status
       2  from user_indexes where status <> 'VALID';

INDEX_NAME                                TABLE_NAME                                STATUS
-----
EMP_EMAIL_UK                              EMPLOYEES                                  UNUSABLE
EMP_EMP_ID_PK                              EMPLOYEES                                  UNUSABLE
EMP_DEPARTMENT_IX                          EMPLOYEES                                  UNUSABLE
EMP_JOB_IX                                EMPLOYEES                                  UNUSABLE
EMP_MANAGER_IX                             EMPLOYEES                                  UNUSABLE
EMP_NAME_IX                                EMPLOYEES                                  UNUSABLE

6 rows selected.



SQL>
```

6. Using Enterprise Manager, reorganize all the indexes in the HR schema that are marked as UNUSABLE.
- a) In Enterprise Manager, on the page displaying the EMP_EMP_ID_PK index, select **Reorganize** in the **Actions** list, and then click **Go**.
- b) Click **Add**, to add each of the other five indexes to the reorganization operation.
- c) In the Add screen, choose Indexes for the **Type** drop-down list, and enter hr in the **Schema** field.
- d) Click **Search**.

Solutions for Practice 13: Performance Management (continued)

- e) Select the five other indexes whose names start with “EMP_.”


Search

Type	Indexes	
Schema	hr	
Object Name		
Partition Name		
Tablespace		

Search

Available Objects: Indexes

Select All | Select None

Select	Name 
<input type="checkbox"/>	HR.DEPT_ID_PK
<input type="checkbox"/>	HR.DEPT_LOCATION_IX
<input checked="" type="checkbox"/>	HR.EMP_DEPARTMENT_IX
<input checked="" type="checkbox"/>	HR.EMP_EMAIL_UK
<input checked="" type="checkbox"/>	HR.EMP_JOB_IX
<input checked="" type="checkbox"/>	HR.EMP_MANAGER_IX
<input checked="" type="checkbox"/>	HR.EMP_NAME_IX
<input type="checkbox"/>	HR.JHIST_DEPARTMENT_IX
<input type="checkbox"/>	HR.JHIST_EMPLOYEE_IX
<input type="checkbox"/>	HR.JHIST_EMP_ID_ST_DATE_PK

- f) Click **OK**.
- g) Click **Next**.
- h) Keep all the default settings for **Options**, and then click **Next**.
- i) Note that there are no problems reported on **Impact Report**, and then click **Next**.
- j) On the Schedule page, enter `oracle` and `oracle` for Username and Password under **Host Credentials**.
- k) Click **Next**.
- l) On the Review page, click **Submit Job**.

Solutions for Practice 13: Performance Management (continued)

- m) After the Confirmation page appears, click the job name to see its status. Click **Reload** on your browser until you see the job has succeeded.

Job Run: REORGANIZE_ORCL.ORACLE.COM_1 at Oct 28, 2005 2:03:27 PM GMT-07:00

Page Refreshed Oct 28, 2005 2:03:46 PM Delete Run Edit

Scheduled **Oct 28, 2005 2:03:27 PM GMT-07:00**

Targets **orcl.oracle.com**

Type **Reorganize**

Owner **DBA1**

Description

Executions

Status All Go

A job run is made up of one or more executions. An execution has zero or more targets. The Retry operation will run immediately and may skip steps that are already completed successfully. The Stop and Suspend operations will wait for the current step of a running execution to complete. The Suspend operation will prevent a scheduled execution from running at its scheduled time. A suspended job can be resumed later.

Stop

Select	Targets	Status	Started	Ended	Elapsed Time (seconds)
	orcl.oracle.com	Succeeded	Oct 28, 2005 2:03:29 PM GMT-07:00	Oct 28, 2005 2:03:45 PM GMT-07:00	16

7. Return to the SQL*Plus session where the HR user is logged in, and run the lab_13_07.sql script to execute the same kind of query. Then, repeat the steps to see the plan of the last SQL statement executed by this session, to see if the plan has changed.

- a) Enter the following at the SQL*Plus prompt:

```
SQL> @lab_13_07.sql
```

- b) Repeat the tasks listed in step 3. Note that the plan now uses the index.

Operation	Object	Object Type
▼ SELECT STATEMENT		
▼ TABLE ACCESS BY INDEX ROWID	EMPLOYEES	TABLE
INDEX UNIQUE SCAN	EMP EMP ID PK	INDEX (UNIQUE)

- c) Quit the SQL*Plus session.
8. What is the difference in execution plans, and why?

Answer: The statement execution uses a unique index scan instead of a full table scan, because the index is usable after your index reorganization.

Solutions for Practice 13: Performance Management (continued)

9. Simulate a working load on your instance by running the `lab_13_09.sql` script as the DBA1 user. Please note the SID value for task 10.

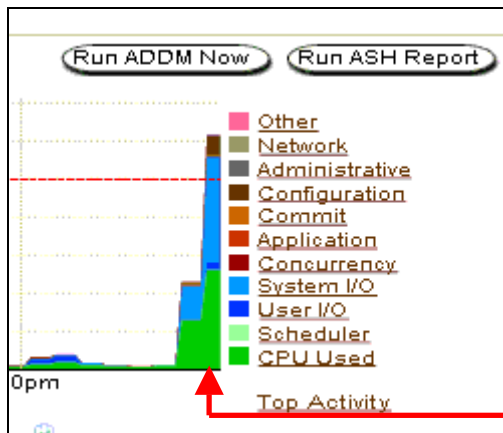
Possible answer: 147 (Your answer is most likely a different one).

This script takes about 20 minutes to complete. So, run it in a separate terminal window and continue with this practice exercise while it runs.

Note: Because this script generates a fairly heavy load in terms of CPU and disk I/O, you will notice that response time for Database Control is slower.

```
$ sqlplus dba1/oracle
SQL> @lab_13_09.sql
```

- a) In Enterprise Manager, navigate to the Performance page, and investigate system performance.
- b) You may need to wait a minute or two to see the effects of the load generation script appear on the graphs.



Wait to see the beginning of a spike in the **Average Active Sessions** graph before proceeding.


Question 1: In the **Average Active Sessions** graph, which are the two main categories that active sessions are waiting for?

Answer: System I/O and CPU Used

Solutions for Practice 13: Performance Management (continued)

Question 2: In the Configuration category of waits, what is one of the contributors to the wait time? Click Configuration to see the graph.

Answer: Any one of these:



Question 3: Click Back, and then click Physical Writes on the Instance Disk I/O graph. Determine which process is doing the most writing to the disk.

Answer: DBW0

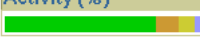
- c) Click **Back**.
- d) Click **Top Activity** in the **Additional Monitoring Links** region.
- e) Click the **SQL ID** of the first SQL statement listed in the **Top SQL** region.
- f) See the first SQL statement.

Question 4: What SQL statement is causing the most waits?

Answer: delete from sh.sales_copy

- 10. Kill the session that is generating the load. Use the session ID recorded in step 9. The session ID should be listed at the top of the list on the right side of the page, under the **Top Sessions** region.


- a) Click the SID number for the session ID recorded earlier. This is found under the heading **Detail for Selected 5 Minute Interval**.

Detail for Selected 5 Minute Interval					
Start Time Nov 2, 2005 4:33:50 PM					
Activity (%)	SID	User	Program	Service	Plan Hash Value
 100.00	131	DBA1	sqlplus@EDRSR4P1 (TNS V1-V3)	SYS\$USERS	2494303166
Statistics	Activity	Plan	Tuning Information		


Solutions for Practice 13: Performance Management (continued)

- b) On the Session Details page, click **Kill Session**, and then click **Yes** to confirm.

Note: If you remain on this Session Details page long enough for a few automatic refreshes to be done, you may see the following warning, which means you are attempting to refresh information about a session that's already been killed. You can ignore this warning.



- c) Click **Top Activity** in the navigation history at the top of the page. Note that the session activity in the database has declined considerably.



Solutions for Practice 14: Backup and Recovery Concepts

Background: Your ORCL database is ready to move from test or development into production. Configure your database to reduce the chances of failure or data loss.

1. Verify that you have two control files to ensure redundancy.
 - a) Invoke Enterprise Manager as the DBA1 user in the SYSDBA role for your ORCL database.
 - b) Select **Administration** > Database Administration > Storage > **Control Files**.

Control Files

[General](#) [Advanced](#) [Record Section](#)

Backup To Trace

Control File Mirror Images

Oracle strongly recommends that your database has a minimum of two control files and that they are located on separate disks. If a control file is damaged due to a disk failure, it could be restored using the intact copy of the control file from the other disk. You can specify their location in the database's initialization parameter file.

Valid	File Name	File Directory
VALID	o1_mf_1p4sbn1o_.ctl	/u01/app/oracle/oradata/ORCL/controlfile/
VALID	o1_mf_1p4sbncs_.ctl	/u01/app/oracle/flash_recovery_area/ORCL/controlfile/

[General](#) [Advanced](#) [Record Section](#)

- c) *Question 1:* How would you add another control file if you needed to?

Answer: Adding a control file is a manual operation. To perform this, you must:

- Shut down the database
- Use the operating system to copy an existing control file to the location where you want your new file to be
- Start the database by using Enterprise Manager. Unlike a normal startup, you would use Advanced Options to select a different startup mode. Select “Start the instance” to leave the instance in the NOMOUNT state.
- Edit the CONTROL_FILES initialization parameter to point to the new control file
- Continue the STARTUP database operation until the database is in an open state

Note: This answer does not apply to an OMF database, as the control files in that case would have to all be recreated.





2. Check how many members each redo log group has. Ensure that there are at least two redo log members in each group. In what directory or directories are the redo log files stored?

Solutions for Practice 14: Backup and Recovery Concepts (continued)

- a) Select **Administration** > Database Administration > Storage > **Redo Log Groups**, and note how many members are in the “# of Members” column. There should be two per group.

Selection Mode Single Create

Edit View Delete Actions Clear logfile Go

Select	Group 	Status	# of Members	Archived	Size (KB)	Sequence	First Change#
	<u>1</u>	Inactive	2	No	51200	26	1408043
	<u>2</u>	Inactive	2	No	51200	27	1424729
	<u>3</u>	Current	2	No	51200	28	1452632

- b) Select the first group, and then click **Edit** to see the member file names. Note that one member is in directories under the `oradata` directory, and another is in the Flash Recovery Area. Click the browser’s **Back** button.

Redo Log Members			Add
			Edit Remove
Select	File Name	File Directory	
	o1_mf_1_1p4sbq1m_.log	/u01/app/oracle/oradata/ORCL/onlineolog/	
	o1_mf_1_1p4sbs8b_.log	/u01/app/oracle/flash_recovery_area/ORCL/onlineolog/	

Note: In a production database, you want to ensure that the two members are on different hard drives, preferably with different disk controllers, to minimize the risk of any single hardware failure destroying an entire log group.

3. You notice that, for each log group, the Archived column has a value of No. This means that your database is not retaining copies of redo logs to use for database recovery, and in the event of a failure, you will lose all data since your last backup. Place your database in ARCHIVELOG mode, so that redo logs will be archived.

- a) Create a new directory `/u01/app/oracle/archive` as the destination for the redo log files.

- In a terminal window, enter `cd /u01/app/oracle`.

- Then, enter `mkdir archive`.

- Optionally, enter `ls -l` to view your new directory and its OS permissions.

Solutions for Practice 14: Backup and Recovery Concepts (continued)

- b) In Enterprise Manager, select **Maintenance** > High Availability > Backup/Recovery Settings > **Recovery Settings**.
 - c) In the **Media Recovery** region, select the **ARCHIVELOG Mode** check box.
 - d) Verify that Log Archive Filename Format contains %t, %s, and %r.
4. Configure redundant archive log destinations—one to the Flash Recovery Area and the other to /u01/app/oracle/archive.

Media Recovery

The database is currently in NOARCHIVELOG mode. In ARCHIVELOG mode, hot backups and recovery to the latest time is possible, but you must provide space for logs. If you change the database to ARCHIVELOG mode, you should make a backup immediately. In NOARCHIVELOG mode, you can make only cold backups and data may be lost in the event of database corruption.

☒ ARCHIVELOG Mode*

Log Archive Filename Format*
The naming convention for the archived log files. %s: log sequence number; %t: thread number; %S and %T: padding the filename to the left with zeroes.

Number	Archive Log Destination	Quota (512B)	Status	Type
1	/u01/app/oracle/archive/			Local
2	<input type="text" value="Archive Log Destination"/>			Local
3	<input type="text"/>			Local
4	<input type="text"/>			Local
5	<input type="text"/>			Local
6	<input type="text"/>			Local
7	<input type="text"/>			Local
8	<input type="text"/>			Local
9	<input type="text"/>			Local
10	USE_DB_RECOVERY_FILE_DEST	n/a	VALID	Local

☒ **TIP** It is recommended that archive log files be written to multiple locations spread across the different disks.
☒ **TIP** You can specify up to 10 archive log destinations.

Flash Recovery Area

It is highly recommended that you use flash recovery area to automate your disk backup management.

Flash Recovery Area Location

Notice that the database is preconfigured to save archived logs to the Flash Recovery Area by default (Archive Log Destination 10). Add an additional Archive Log Destination so that you will have redundant copies of your log files.

- a) In **Archive Log Destination** number 1, enter /u01/app/oracle/archive/. The directory path should end with a slash. Leave **Quota** blank.

Note: You must create the directory, if it does not already exist. You already did this in step (3).

Solutions for Practice 14: Backup and Recovery Concepts (continued)

- b) Optionally, click **Show SQL**, review the statements, and then click **Return**.

Show SQL

```
ALTER SYSTEM SET log_archive_dest_1 =  
"LOCATION=/u01/app/oracle/archive/ OPTIONAL REOPEN=300" SCOPE=BOTH  
ALTER SYSTEM SET log_archive_dest_10 =  
"LOCATION=USE_DB_RECOVERY_FILE_DEST OPTIONAL REOPEN=300" SCOPE=BOTH
```

- c) Click **Apply**. When prompted whether you want to restart the database now, click **Yes**.

Database: orcl > Restart Database:Specify Host and Target Database Credentials

Logged in As SYS

Restart Database:Specify Host and Target Database Credentials

Specify the following credentials in order to restart the database.

Host Credentials

Specify the OS user name and password to login to target database machine.

* Username

* Password

Database Credentials

Specify the credentials for the target database.
To use OS authentication, leave the user name and password fields blank.

* Username

* Password

Database **orcl**

* Connect

As

☒ Save as Preferred Credential

Note that you need to login to the database as SYSDBA or SYSOPER in order to restart the database.

- d) Enter the credentials to restart the database, and then click **OK**.

- e) When asked to confirm, click **Yes** again.

Now that your database is in ARCHIVELOG mode, it will continually archive a copy of each online redo log file before reusing it for additional redo data.

Note: Remember that this consumes space on the disk and that you must regularly back up older archive logs to some other storage.

Solutions for Practice 15: Performing Database Backups

Background: Your database is ready to move from development and test into production. Ensure that your database is configured so that recovery is possible without loss of data.

1. What is the difference between a backup set and an image copy?

Answer: A backup set contains data and archive log files packed in an Oracle proprietary format. Files must be extracted before use. Image copies are the equivalent of operating system file copies and can be used to restore operations immediately.

2. What is the destination of any disk backups that are done?
 - a) In Enterprise Manager, select **Maintenance** > High Availability > Backup/Recovery Settings > **Backup Settings**. Note the message under the Disk Backup Location that says the Flash Recovery Area is the current disk backup location.

Backup Settings

Device | Backup Set | Policy

Disk Settings

Parallelism Test Disk Backup

Concurrent streams to disk drives

Disk Backup Location

Flash recovery area is your current the disk backup location. If you would like to override the disk backup location, specify an existing directory or diskgroup name.

3. Test making a backup to disk, as a backup set, with `oracle` for Host Credentials.
 - a) Select **Backup Set** as your Disk Backup Type.
 - b) Scroll to the bottom and enter `oracle` and `oracle` for **Host Credentials Username** and **Password** for your server.
 - c) Click **Test Disk Backup**.
 - d) When the test finishes, click **OK**.

Solutions for Practice 15: Performing Database Backups (continued)

4. Back up your entire database, without archived logs, while the database is open for user activity. This backup should be the base for an incremental backup strategy.

Question: What prerequisite *must* be met to create a valid backup of a database without shutting it down?

Answer: The database must be in ARCHIVELOG mode. Backups made with the database open but not in ARCHIVELOG mode cannot be used for recovery.

- a) Select **Maintenance** > High Availability > Backup/Recovery > **Schedule Backup**.
- b) Select **Whole Database** as the object to be backed up.
- c) Enter `oracle` and `oracle` for **Host Credentials Username** and **Password** for your server.

The screenshot shows two panels from the Oracle Enterprise Manager interface. The top panel, titled 'Customized Backup', has a heading 'Select the object(s) you want to back up.' and a 'Schedule Customized Backup' button. It lists five options: 'Whole Database' (selected with a radio button), 'Tablespaces', 'Datafiles', 'Archivelogs', and 'All Recovery Files on Disk' (with a note: 'These files include all archivelogs and disk backups that are not already backed up to tape'). The bottom panel, titled 'Host Credentials', has a heading 'To perform a backup, supply operating system login credentials to access the target database.' and contains fields for 'Username' (filled with 'oracle') and 'Password' (filled with '*****'). There is also a checkbox labeled 'Save as Preferred Credential' which is currently unchecked.

- d) Click **Schedule Customized Backup**.
- e) Select **Full Backup** for your Backup Type, and select the **Use as the base of an incremental backup strategy** check box.
- f) Select **Online Backup** as Backup Mode.
- g) In the **Advanced** region, deselect the **Also backup all archived logs on disk** check box, and then click **Next** to continue.

Solutions for Practice 15: Performing Database Backups (continued)

- h) Select **Disk** for your backup location (notice that your Disk Backup Location is retained from step [2]). Click **Next** to continue.
- i) Accept all the defaults on the Schedule page, and then click **Next** to continue.

Job

* Job Name

BACKUP_ORCL.Oracle.com_01

Job Description

Whole Database Backup

Schedule

Time Zone GMT -7:00

Start

☒ Immediately

☐ Later

Date

Jun 9, 2005

(example: Jun 9, 2005)

Time

2

00

☒ AM ☐ PM

Repeat

☒ One Time Only

☐ Interval

☐ Monthly

☐ Yearly

Frequency

1

Minutes

Repeat Until

☒ Indefinite

☐ Custom

Date

Jun 9, 2005

(example: Jun 9, 2005)

Time

3

30

☒ AM ☐ PM

(Ignored except when repeating by minutes or hours.)

- j) Click **Submit Job** to perform the online database backup.

Solutions for Practice 15: Performing Database Backups (continued)

- k) Click **View Job** to monitor the status of the backup job. This backup takes approximately five minutes to complete.

Execution: orcl.oracle.com

Page Refreshed Jun 9, 2005 4:34:54 AM Delete Run Edit

Summary

The Stop and Suspend operations will wait for the current step to complete. A suspended job can be resumed later, at the next step. Stop

Status	Running	Type	Database Backup
Scheduled	Jun 9, 2005 4:34:35 AM GMT-07:00	Owner	SYS
Started	Jun 9, 2005 4:34:37 AM GMT-07:00	Description	Whole Database Backup
Start Delayed	2 seconds	Host Username	oracle
Elapsed Time	16 seconds	Database	(DESCRIPTION=(ADDRESS LIST=(ADDR...
		Connect String	
		Database Username	SYS
		Database Role	[SYSDBA]
		Oracle Home	[u01/app/oracle/product/10.2.0/...]
		Oracle SID	[orcl]
		Version 10g or higher	YES
		Backup Strategy	advanced
		Offline Backup	NO
		Blackout	NO
		Database Name	ORCL
		Backup Script	Show

5. Schedule nightly disk-based incremental online backups for your whole database, without archived logs. Schedule it for 11:00 p.m. The schedule should be in effect indefinitely.
- In Enterprise Manager, select **Maintenance** > High Availability > Backup/Recovery > **Schedule Backup**.
 - Select **Whole Database** as the object to be backed up.
 - Enter `oracle` and `oracle` for **Host Credentials Username** and **Password** for your server, and then click **Schedule Customized Backup**.
 - Select **Incremental Backup (Level 1)** for your Backup Type.
 - Select **Online Backup** as Backup Mode.
 - In the **Advanced** region, deselect the **Also backup all archived logs on disk** check box, and then click **Next** to continue.
 - Select **Disk** as your backup location, and then click **Next** to continue.

Solutions for Practice 15: Performing Database Backups (continued)

- h) Change **Job Name** to `Nightly_Backup` and accept the default value for **Job Description**.
- i) Select **Later** in the **Start** region. Accept today's date and use the drop-down lists and option buttons to select 11:00 p.m. for **Time**.
- j) In the **Repeat** area, select **Interval**, and **Frequency 1 Days**.
- k) Select **Indefinite** in the **Repeat Until** region, and then click **Next**.

Schedule Customized Backup: Schedule

Database **orcl.oracle.com**
Backup Strategy **Customized Backup**
Object Type **Whole Database**

Cancel Back Step 3 of 4 Next

Job

* Job Name
Job Description

Schedule

Time Zone

Start

☐ Immediately
☒ Later
Date
(example: Jun 9, 2005)
Time ☐ AM ☒ PM

Repeat

☐ One Time Only
☒ Interval
Frequency
☐ Monthly
☐ Yearly

Repeat Until

☒ Indefinite
☐ Custom
Date
(example: Jun 9, 2005)
Time ☐ AM ☐ PM
(Ignored except when repeating by minutes or hours.)

- l) Click **Submit Job**, and then click **OK**.
- m) Navigate to **Maintenance > Related Links > Jobs** to see the scheduled job in the **Job Activity** list.

Solutions for Practice 16: Performing Database Recovery

Background: Many failures of the Oracle database can be traced to some sort of media failure, such as disk or controller failure. Recover your database from a variety of simulated media failures. SQL script files are provided for you in the `/home/oracle/labs` directory. If needed, use appendix C for Linux and appendix D for SQL syntax. Note that where OS file names are mentioned, your system may possibly have different file names than shown here.

1. Recover from the loss of a control file.

- a) As the DBA1 user, run the `lab_16_01_a.sql` script to prepare some procedures to be called by the rest of this practice.

```
@$HOME/labs/lab_16_01_a.sql
```

- b) Now run the `lab_16_01_b.sql` script. This script deletes one of your control files.

```
@$HOME/labs/lab_16_01_b.sql
```

- c) The Help desk begins receiving calls saying that the database appears to be down. Troubleshoot and recover as necessary. Use Enterprise Manager's Database page to attempt to start up the database, and use SQL*Plus, if needed.
- In Enterprise Manager, navigate to the Database page. It reports that the database is down and offers you the chance to start it up again. Attempt to do so by clicking **Startup**. You may need to wait one or two minutes before the page appears as described. If you see a Connection Refused message, ignore it; the connection will eventually be established.
 - Enter host credentials as `oracle` for **Username** and **Password**, and then click **OK**.
 - When asked whether you are sure that you want to start the database, click **Yes**.

Solutions for Practice 16: Performing Database Recovery (continued)

- d) The startup of the instance fails with Enterprise Manager, and you can get no other information to explain the failure. So use the command-line tools.

- Connect to the instance with SQL*Plus as sysdba and check the current status of the instance.

```
sqlplus / as sysdba
select status from v$instance;
```

- The instance status is STARTED, which means that the database is in the NOMOUNT stage. Attempt to mount the database by entering this:

```
SQL> alter database mount;
```

```
SQL> alter database mount;
alter database mount
*
ERROR at line 1:
ORA-00205: error in identifying control file, check alert log for more info
```

- e) The instance cannot move to the mount stage because it cannot find one of the control files. Check the last 10 rows of the alert log to see which control file is the problem.

```
SQL> host tail -10 $ORACLE_BASE/admin/orcl/bdump/aler*
```

Output:

```
SQL> host tail -10 $ORACLE_BASE/admin/orcl/bdump/aler*
ORA-205 signalled during: ALTER DATABASE MOUNT...
Sun Oct 30 13:35:01 2005
alter database mount
Sun Oct 30 13:35:01 2005
ORA-00202: control file:
'/u01/app/oracle/flash_recovery_area/ORCL/controlfile/o1_mf_1p4sbncs_.ct
l'
ORA-27037: unable to obtain file status
Linux Error: 2: No such file or directory
Additional information: 3
Sun Oct 30 13:35:01 2005
ORA-205 signalled during: alter database mount...
```

Solutions for Practice 16: Performing Database Recovery (continued)

- f) The control file in the Flash Recovery Area is missing. Restore the missing control file by copying the existing control file, and then mount and open the database.

```
SQL> host
$ cd /u01/app/oracle/oradata/ORCL/controlfile
$ ls -l
total 6908
-rw-r----- 1 oracle oinstall 7061504 Oct 30 13:32 o1_mf_1p4sbn1o_.ctl
$ cp o1_mf_1p4sbn1o_.ctl
/u01/app/oracle/flash_recovery_area/ORCL/controlfile/o1_mf_1p4sbncs_.ctl
$ exit
```

```
SQL> alter database mount;
```

```
SQL> alter database open;
```

- g) Why did you have to use two commands to move the instance state from NOMOUNT to OPEN?

Answer: Because the ALTER DATABASE command enables you to change only one state level per command

- h) Why did you use operating system commands to restore the control file instead of using Oracle Recovery Manager?

Answer: Because all control files are identical. As long as any one control file is intact, it can be used to restore the others.

Solutions for Practice 16: Performing Database Recovery (continued)

2. Recover from the loss of an application data file.

- a) Start a SQL*Plus session as the DBA1 user, and run the lab_16_02.sql script. This script deletes one of your application data files.

```
$ sqlplus dba1/oracle @$HOME/labs/lab_16_02.sql
```

- b) The Help desk has received a call from a user who is unable to access the COUNTRIES table in the HR application schema. Count the rows in the table to see whether there is a problem.

```
SQL> select count(*) from HR.COUNTRIES;
select count(*) from HR.COUNTRIES
                        *
ERROR at line 1:
ORA-01116: error in opening database file 5
ORA-01110: data file 5:
'/u01/app/oracle/oradata/ORCL/datafile/o1_mf_example_1p4sd3y2_.dbf'
ORA-27041: unable to open file
Linux Error: 2: No such file or directory
Additional information: 3
```

- c) Troubleshoot and recover as necessary. The error message suggests that the data file for the EXAMPLES tablespace is corrupt or missing. Using operating system commands, verify that there is a problem with the file.

```
SQL> host ls
/u01/app/oracle/oradata/ORCL/datafile/o1_mf_example_1p4sd3y2_.dbf
ls: /u01/app/oracle/oradata/ORCL/datafile/o1_mf_example_1p4sd3y2_.dbf:
No such file or directory
```

- d) Recover the data file to the current time, specifying the missing data file to be recovered.

- In Enterprise Manager, select **Maintenance** > High Availability > Backup/Recovery > **Perform Recovery**.
- Select **Datafiles** from the **Object Type** drop-down list.
- In the Object Level Recovery region, select **Recover to current time** for the **Operation Type**.
- Enter the host credentials as **oracle** and **oracle** for **Username** and **Password**, and then click **Perform Object Level Recovery**.

Solutions for Practice 16: Performing Database Recovery (continued)

- On the Datafiles page, select the data file in question.

Select	Datafile Name	Datafile Number	Status	Enabled
<input checked="" type="checkbox"/>	/u01/app/oracle/oradata/ORCL/datafile/o1_mf_example_1p4sd3y2_.dbf	5	ONLINE	READ WRITE

- Click **Next**.
- Because the problem is simply a deleted file rather than a bad hard drive, there is no need to restore the file to a different location. Select **No. Restore the files to the default location**, and then click **Next**.
- Click **Submit**. (It will take one or two minutes for the operation to complete.)

e) When you see the Operation Succeeded message, ensure that the restored data file is online.

- In Enterprise Manager, select **Administration** > Database Administration > Storage > **Datafiles**.

Datafiles

Object Type: Datafile

Search
Select an object type and optionally enter an object name to filter the data that is displayed in your results set.
Object Name:

By default Datafiles are case-sensitive searches. To run an exact match, double quote the search string. You can use the wildcard symbol (%) in a double quoted string.

Select	File Name	Tablespace	Status	Size (MB)	Used (MB)	Used (%)
	/u01/app/oracle/oradata/ORCL/datafile/inventory01.dbf	INVENTORY	ONLINE	50.000	42.188	84.38
	/u01/app/oracle/oradata/ORCL/datafile/o1_mf_example_1pbkt0j1_.dbf	EXAMPLE	ONLINE	100.000	68.250	68.25
	/u01/app/oracle/oradata/ORCL/datafile/o1_mf_sysaux_1p4s741n_.dbf	SYSAUX	ONLINE	270.000	264.812	98.08
	/u01/app/oracle/oradata/ORCL/datafile/o1_mf_system_1p4s740z_.dbf	SYSTEM	SYSTEM	490.000	480.688	

f) Verify that the COUNTRIES table is now accessible.

```
SQL> select count(*) from HR.COUNTRIES;

COUNT (*)
-----
         25
```

Solutions for Practice 16: Performing Database Recovery (continued)

3. Recover from the loss of a system data file.

- a) Why is recovery from the loss of a system data file or a data file belonging to an undo tablespace different from recovering an application data file?

Answer: Because recovery of system or undo data files must be done with the database closed, whereas recovery of an application data file can be done with the database open and available to users.

- b) As SYSDBA, run the lab_16_03.sql script. This script deletes the system data file.

```
SQL> @lab_16_03.sql

PL/SQL procedure successfully completed.

ORACLE instance shut down.
SQL>
```

- c) In Enterprise Manager, review the Database home page. The database is shut down, so you click **Startup** to try to open it. If you see a message that says the connection was refused, dismiss it, and reenter the EM home page URL in the browser.

- Enter the host credentials as `oracle` and `oracle` for the host **Username** and **Password**, and enter `DBA1` and `oracle` for the database credentials, and then click **OK**.



- Click **Yes**.

Solutions for Practice 16: Performing Database Recovery (continued)

- d) This command will fail with the database left in the MOUNT state, because there is a data file missing from the SYSTEM tablespace.



- e) Click **Perform Recovery**.
- If prompted, enter host (oracle/oracle) and database (dba1/oracle AS SYSDBA) credentials, and then click **Continue**.
 - In the **Object Level Recovery** region, select **Datafiles** for Object Type.
 - Select **Recover to current time** for **Operation Type**.
 - Fill in the host credentials if not already set, and then click **Perform Object Level Recovery**.
 - Select the data file for the SYSTEM tablespace, and then click **Next**.
 - Because the problem is simply a deleted file rather than a bad hard drive, there is no need to restore to a different location. Select **No. Restore the files to the default location**, and then click **Next**.
 - Click **Submit**. It will take three to four minutes for the operation to complete.
- f) When you see the Operation Succeeded message, click **Open Database**.
- g) After you see the success message, click **OK**, and then verify that the database is open and operating normally by logging into EM as DBA1/oracle, as SYSDBA, and reviewing the Database home page.

Solutions for Practice 17: Performing Flashback

Background: You decide to gain hands-on experience in some of the flashback functionality. To avoid impacting other users, you will first copy the DEPARTMENTS table of the HR schema to DEPARTMENTS2.

1. Log in to SQL*Plus as DBA1 user and create a new HR.DEPARTMENTS2 table based on the HR.DEPARTMENTS table.

a) Log in to SQL*Plus as dba1/oracle.

```
$ sqlplus dba1/oracle
```

b) Enter the following command to create the copy table:

```
SQL> create table hr.departments2 as select * from hr.departments;
```

c) Count the rows in the DEPARTMENTS2 table. There should be 27 rows.

```
SQL> SELECT COUNT(*) FROM HR.DEPARTMENTS2;

COUNT(*)
-----
        27
```

2. Drop the HR.DEPARTMENTS2 table, and then verify that it has indeed been dropped.

```
SQL> DROP TABLE HR.DEPARTMENTS2;

Table dropped.

SQL> SELECT * FROM HR.DEPARTMENTS2;
SELECT * FROM HR.DEPARTMENTS2
                *
ERROR at line 1:
ORA-00942: table or view does not exist
```

3. Use the FLASHBACK TABLE command to restore the table. Count the rows in the DEPARTMENTS2 table.

```
SQL> FLASHBACK TABLE hr.departments2 TO BEFORE DROP;

Flashback complete.

SQL> SELECT COUNT(*) FROM hr.departments2;

COUNT(*)
-----
        27
```

Solutions for Practice 17: Performing Flashback (continued)

4. Run the `lab_17_04.sql` script to insert three rows into the `HR.DEPARTMENTS2` table by using three separate transactions. The new rows have `DEPARTMENT_ID` values of 280, 290, and 300.

```
SQL> @lab_17_04.sql
```

5. Use Enterprise Manager to perform flashback to the version of the table where only the first of the three new rows is present (with `DEPARTMENT_ID = 280`). First, evaluate row changes to decide on a point in time. Limit your analysis to the new rows just added: where `department_id >= 280`. If you receive an error while performing the flashback, you may need to enable row movement on the table. See the next step.
- a) In Enterprise Manager, select **Maintenance** > High Availability > Backup/Recovery > **Perform Recovery**.
 - b) Select Tables from the **Object Type** drop-down list, and then select **Flashback Existing Tables** for **Operation Type**. Click **Perform Object Level Recovery**.

Perform Recovery

Whole Database Recovery

☒ Recover to the current time or a previous point-in-time
Datafiles will be restored from the latest usable backup as required.

☐ Restore all datafiles
Specify Time, SCN or log sequence. The backup taken at or prior to that time will be used. No recovery will be performed in this operation.

☐ Recover from previously restored datafiles

Object Level Recovery

Object Type Tables ▾

Operation Type ☒ Flashback Existing Tables
☐ Flashback Dropped Tables

Host Credentials

To perform recovery, supply operating system login credentials to access the target database.

* Username

* Password

☐ Save as Preferred Credential

- c) Select **Evaluate row changes and transactions to decide on a point in time**, and enter `HR.DEPARTMENTS2` as the fully qualified name of the table in the **Table** field, and then click **Next**.

Solutions for Practice 17: Performing Flashback (continued)

- d) Highlight DEPARTMENT_ID under **Available Columns**, and then click the **Move** button to move it under **Selected Columns**. Under Step 2, enter a WHERE clause that will select the added rows. For example, earlier you added rows with DEPARTMENT_IDs of 280, 290, and 300, so the “WHERE department_id >= 280” clause is suitable.

Step 1. Choose Columns

Available Columns

DEPARTMENT_NAME
MANAGER_ID
LOCATION_ID

>

Move

>>

Move All

<

Remove

<<

Remove All

Selected Columns

DEPARTMENT ID

Step 2. Bind The Row Value

Specify a where clause based on the columns selected above to narrow the s

where department_id >= 280

- e) Click **Next** to continue.
- f) You can now review the rows under Flashback Versions Query Result.

Flashback Versions Query Result					
Select	Flashback SCN	Flashback Timestamp	Transaction ID	Operation	DEPARTMENT_ID
	6800080	Jun 9, 2005 5:15:32 PM	0A002B00A6030000	INSERT	300
	6800061	Jun 9, 2005 5:15:15 PM	03002400CB030000	INSERT	290
	6800049	Jun 9, 2005 5:14:53 PM	05000600FA030000	INSERT	280

- g) Under **Flashback Versions Query Result**, select the middle transaction from the list to flash back. Click **Next** to continue.

Solutions for Practice 17: Performing Flashback (continued)

- h) On the next page, you are asked to include any tables related to or dependent on the table that you are flashing back. Because there are none, click **Next** to continue.

Perform Object Level Recovery: Flashback Tables

Object Type	Tables	<input type="button" value="Cancel"/> <input type="button" value="Back"/> Step 4 of 7 <input type="button" value="Next"/>
Operation Type	Flashback Existing Tables	

Your application may have tables that are logically related to this table. Specify all such tables that must be flashed back to the SCN you selected.

Evaluated Table Name	hr.departments2
Flashback Time	Jun 9, 2005 05:15 PM
Flashback SCN	6800061


HR.DEPARTMENTS2

- i) Review the SQL statement that you are about to execute. Click **Show SQL**, view the SQL statement, click **OK**, and then click **Submit**.

```
FLASHBACK TABLE HR.DEPARTMENTS2 TO SCN 6800061
```

6. On the next page, you find that the operation has failed because row movement is not enabled for the table. You may recall from the lesson that row movement must be enabled for this feature to work.

- a) Note this error:



Error

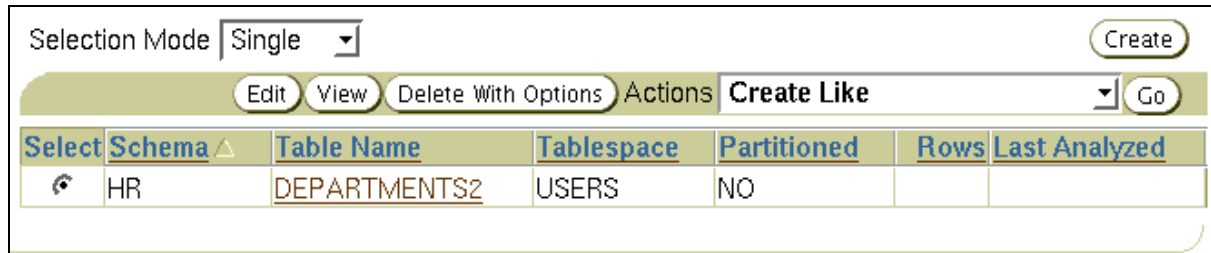
Examine and correct the following errors then re-try the operation:

Error - ORA-08189: cannot flashback the table because row movement is not enabled

- b) To enable row movement for this table, select **Administration** > Schema > Database Objects > **Tables**.
- c) Enter HR in the **Schema** field and DEPARTMENTS2 in the **Object Name** field, and then click **Go**.

Solutions for Practice 17: Performing Flashback (continued)

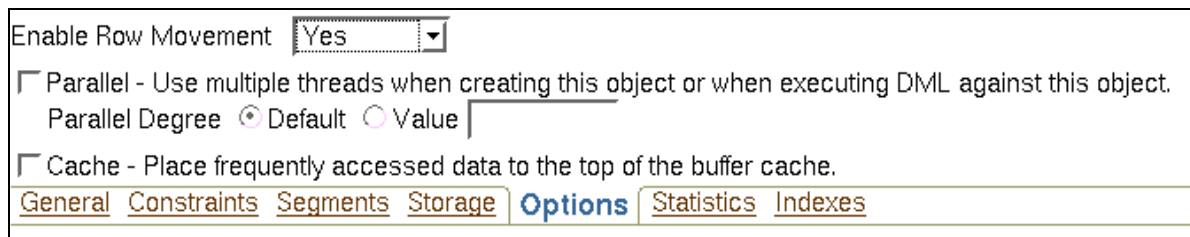
- d) One table is displayed in the **Results** region and is selected by default. Click **Edit**.



Select	Schema	Table Name	Tablespace	Partitioned	Rows	Last Analyzed
<input checked="" type="radio"/>	HR	DEPARTMENTS2	USERS	NO		

The Edit Table page appears.

- e) Click the **Options** tab.
- f) From the **Enable Row Movement** drop-down list, select **Yes**.



Enable Row Movement: **Yes**

☐ Parallel - Use multiple threads when creating this object or when executing DML against this object.
Parallel Degree: ☒ Default ☐ Value

☐ Cache - Place frequently accessed data to the top of the buffer cache.

[General](#) [Constraints](#) [Segments](#) [Storage](#) **[Options](#)** [Statistics](#) [Indexes](#)

- g) Click **Show SQL**, and then click **Return**.

```
ALTER TABLE "HR"."DEPARTMENTS2" ENABLE ROW MOVEMENT
```

- h) Click **Apply** to apply the change.

After the table modification is confirmed, you can perform the flashback operation.

Repeat step 5 to perform the version flashback of the table. Because row movement has been enabled, it should succeed this time.



Confirmation

The selected tables, HR.DEPARTMENTS2, have been flashed back.

Solutions for Practice 17: Performing Flashback (continued)

- i) In SQL*Plus, count the rows of the HR.DEPARTMENTS2 table to confirm the flashback operation. Note that there is only one additional row now, not three. Then display the rows where DEPARTMENT_ID >= 280. Note that only one of the original three is remaining.

```
SQL> select count(*) from hr.departments2;

COUNT(*)
-----
         28

SQL> select * from hr.departments2 where department_id >= 280;

DEPARTMENT_ID DEPARTMENT_NAME          MANAGER_ID LOCATION_ID
-----
          280 DUMMY1
```

SQL>

Solutions for Practice 18: Moving Data

Background: In the recent past, you received a number of questions about the HR schema. To analyze them, without interfering in the daily activities, you decide to use the Data Pump Wizard to export the HR schema to file. When you perform the export, you are not sure into which database you will be importing this schema.

In the end, you find out that the only database for which the management approves an import for, is the ORCL database. So you perform the import with the Data Pump Wizard, remapping the HR schema to a newly created HR_TEST schema in the HR_TEST tablespace. To follow best practice guidelines, you also decide to create a DP user who will be a DBA performing Data Pump jobs. For your convenience in class, the creation of the DP user is included in the lab_18_01.sql script. SQL scripts are in the /home/oracle/labs directory. However, there is one step missing.

Then, you receive two data load requests for which you decide to use SQL*Loader.

1. Review the lab_18_01.sql script, which creates the HR_TEST tablespace, the HR_TEST_ROLE role, and the HR_TEST and DP users.
 - a) Double-click the **oracle's Home** icon on your desktop, double-click **labs**, and then double-click **lab_18_01.sql** to review this script.
 - b) Note the passwords for these users.

HR_TEST password: **hr_4test**

DP password: **dp_4test**
 - c) Which additional step do you need to perform to allow the DP user access to Enterprise Manager as Administrator?

You need to log in to Enterprise Manager as the SYS user in the SYDBA role and make the DP user an EM Administrator (as you learned in the lesson titled "Managing the Oracle Instance").
 - d) Close the lab_18_01.sql window.
2. Execute the lab_18_02.sh script. Then, perform the required step to make the DP user an EM administrator.
 - a) Right-click your desktop and select **Open Terminal**. Then navigate to the /home/oracle/labs directory by entering **cd labs**.
 - b) Enter **./lab_18_02.sh** to create the HR_TEST tablespace, the HR_TEST_ROLE role, and the HR_TEST and DP users.
 - c) Log in to Enterprise Manager as the SYS user in the SYSDBA role.

Solutions for Practice 18: Moving Data (continued)

- d) Make the DP user an EM Administrator (as you learned in the lesson titled “Managing the Oracle Instance”). (Select Setup > Create > Enter Name: DP, Password: dp_4test, Super Administrator, checked > Click Finish > Finish > Logout.)

ORACLE Enterprise Manager 10g Database Control

Setup Preferences Help Logout

Database

Setup

Create Administrator: Properties

* Name: dp

* Password: [masked]

* Confirm Password: [masked]

Email Address: [empty]

Specify one or more e-mail addresses separated by a comma or space.

☒ Super Administrator

Cancel Finish

3. Log in to Enterprise Manager as the DP user in the Normal role and export the HR schema.
- a) Invoke Enterprise Manager as the DP user in the Normal role for your ORCL database. The **Connect As** setting should be Normal.

If this is the first time that the DP user logs in, click **I agree** to accept the licensing agreement.

- b) Select **Maintenance** > Data Movement > Move Row Data > **Export to Export Files**.
- c) Select **Schemas**, enter **oracle** as **Username** and **Password**, select **Save as Preferred Credential**, and then click **Continue**.
- d) On the Export: Schemas page, click **Add**, select the HR schema, and then click the **Select** button.

Schemas Options Files Schedule Review

Export: Schemas

Database orcl.oracle.com

Cancel Finish Step 1 of 5 Next

Add

Remove

Select Schemas ▾

HR

- e) Click **Next**.

Solutions for Practice 18: Moving Data (continued)

- f) On the Export: Options page, select DATA_PUMP_DIR from the **Directory Object** drop-down list, and enter **hrexpl.log** as **Log File**.
- g) Review Advanced Options, but do not change, and then click **Next**.
- h) On the Export: Files page, select DATA_PUMP_DIR from the **Directory Object** drop-down list, enter **HREXP%U.DMP** as **File Name**, and then click **Next**.
- i) On the Export: Schedule page, enter **hrexpl** as **Job Name** and **Export HR schema** as **Description**, accept the immediate job start time, and then click **Next**.
- j) On the Export: Review page, click **Show PL/SQL** and review the PL/SQL that the Export Wizard helped you to create.

Export: Review

Database **orcl.oracle.com** Cancel Back Step 5 of 5 Submit Job

Export Type	Schemas
Statistics type	Estimate optimizer statistics when data is imported
Parallelism	1
Files to Export	DATA_PUMP_DIR HREXP%U.DMP
Log File	DATA_PUMP_DIR hrexpl.log

[Hide PL/SQL](#)

Export PL/SQL

```
begin
  dbms_datapump.add_file(handle => h1, filename => 'hrexpl.log', directory => 'DATA_PUMP_DIR',
filetype => 3);
end;
begin
  dbms_datapump.set_parameter(handle => h1, name => 'KEEP_MASTER', value => 0);
end;
begin
  dbms_datapump.metadata_filter(handle => h1, name => 'SCHEMA_EXPR', value => 'IN(''HR'');
end;
```

- k) Click **Submit Job**.

A processing message appears, and then a success message. If not, resolve any errors, which may have occurred.

- l) When the Job Activity Confirmation page appears, click the HREXP job name, and then monitor the job progress by clicking the browser's **Reload** button.

Note: Please wait, not only for the job to be created, but also for the job to complete execution. (It may take two minutes.)

- 4. As the DP user, import the exported HR schema back into the ORCL database, remapping it to the previously created HR_TEST schema.

Solutions for Practice 18: Moving Data (continued)

- a) Invoke Enterprise Manager as the DP user in the Normal role for your ORCL database.
- b) Select **Maintenance** > Data Movement > Move Row Data > **Import from Export Files**.
- c) On the Import: Files page, select DATA_PUMP_DIR from the **Directory Object** drop-down list, and enter **HREXP%U.DMP** as **File Name**.
- d) Select **Schemas** as the Import Type.
- e) Confirm oracle as your **Host Credentials**, and then click **Continue**.

At this point, the export file is read, to verify the contents. Wait for this to complete.

- f) On the Import: Schemas page, click **Add**, select HR, and then click **Select**.
- g) Click **Next**.
- h) On the Import: Re-Mapping page, click **Add Another Row** under **Re-Map Schemas**. Then select **HR_TEST** as **Destination Schema**.
- i) Click **Add Another Row** under **Re-Map Tablespaces**, enter **HR_TEST** as **Destination Tablespace**.

Import: Re-Mapping

Database **orcl.oracle.com** Cancel Finish Back Step 2 of 5 Next

Re-Map Schemas

You can import data that was in one schema into a different schema.

Remove

Select	Source Schema	Destination Schema
	HR	HR_TEST

Add Another Row

Re-Map Tablespaces

You can import data that was in one tablespace into a different tablespace.

Remove

Select	Source Tablespace	Destination Tablespace
		hr_test

Add Another Row

- j) Click **Next**.

Solutions for Practice 18: Moving Data (continued)

- k) On the Import: Options page, select **DATA_PUMP_DIR** from the **Directory Object** drop-down list, enter **hriimport.log** as **Log File**, review the advanced options, but leave them at their default values, and then click **Next**.
- l) On the Import: Schedule page, enter **hriimp** as **Job Name** and **Import HR schema for test purposes** as **Description**.
- m) Select **Later** as **Start** and enter a time between 2 and 5 minutes from now (to give yourself time for the following steps):
 - Click **Next**.
 - On the Import: Review page, review the PL/SQL that the Data Pump Wizard creates for you, and then click **Submit Job**.
 - After the confirmation that the job was successfully created, note the job name (for example HRIMP), and navigate to **Maintenance > Data Movement > Move Row Data > Monitor Export and Import Jobs**.
 - Click your last job (for example, HRIMP) .
 - On the “Monitor Data Pump job” page, click **Reload** in your browser, when it is time to execute your job.

Monitor Data Pump Job

Page Refreshed Jun 10, 2005 5:25:06 PM OK

Job Name

IMPORT000024

Job Status

EXECUTING Change Job State

Percent Done

0

Percent done is for table data only.

Non-fatal Error Count

0

Log File

Maximum Number of Threads in Export Job

1 Change Parallelism

Objects Currently Being Imported

The number of rows in this table corresponds to the degree of parallelism.

Name	Type	Percent Done
		0

Import Files

File Name	Maximum File Size (MB)	Used (MB)	Used (%)
/u01/app/oracle/product/10.2.0/db_1/rdbms/log/HREXP01.DMP	Auto Extend	0	0.0

Solutions for Practice 18: Moving Data (continued)

- *Optional, as this requires quick action:* Click **Change Job State**.

Change Data Pump Job State

CancelOK

Select the job state to which you want to change. The job state will change immediately after you click OK.

☒ Suspend
The Data Pump job will be suspended; it can later be resumed by again doing a Change Data Pump Job State and using the Resume option.

☐ Stop
The Data Pump job will be stopped and cannot be restarted.

- *Optional (continued):* Click **Suspend** and later **Resume** to halt the job and then to continue it again.

n) In the end, you want to see that your job executed 100% without any errors.

Monitor Data Pump Job

Page Refreshed Jun 10, 2005 5:26:05 PMOK

Job Name	IMPORT000024
Job Status	COMPLETING
Percent Done	100
	Percent done is for table data only.
Non-fatal Error Count	0
Log File	
Maximum Number of Threads in Export Job	1

Objects Currently Being Imported

The number of rows in this table corresponds to the degree of parallelism.

Name	Type	Percent Done
		0

Import Files

File Name	Maximum File Size (MB)	Used (MB)	Used (%)
/u01/app/oracle/product/10.2.0/db_1/rdbms/log/HREXP01.DMP	Auto Extend	0	0.0

File Percent used is for table data only.

o) Click **OK**.

p) Verify that the import succeeded by viewing the log file.

```
$ cat </u01/app/oracle/product/10.2.0/db_1/rdbms/log/hrimport.log
```

Note: You may see an error saying that the hr_test object already exists. This is because that user existed when you did the export, and still exists. This is not a problem.

q) Using SQL*Plus, connect to the database as the HR_TEST user.

```
$ sqlplus hr_test/hr_4test
```

r) Select data from tables in the hr_test schema, for verification of the import.

```
SQL> select * from jobs;
```

Solutions for Practice 18: Moving Data (continued)

5. As the DP user, load data into the PRODUCT_MASTER table by using SQL*Loader via Enterprise Manager Database Control.
 - a) Invoke Enterprise Manager as the DP user in the Normal role for your ORCL database.
 - b) Select **Maintenance** > Data Movement > Move Row Data > **Load Data from User Files**.
 - c) Click **Use Existing Control File**.
 - d) If you have not done so before, enter **oracle** as **Username** and as **Password**, click **Save as Preferred Credential**, and then click **Continue**.
 - e) On the Load Data: Control File page, enter **/home/oracle/labs/lab_18_05.ctl** as control file name and path, or use the flashlight icon to select this control file. Click **Next**.
 - f) On the Load Data: Data File page, click **Provide the full path and name on the database server machine** and enter **/home/oracle/labs/lab_18_05.dat** as data file name and path, or use the flashlight icon to select this data file. Click **Next**.
 - g) On the Load Data: Load Method page, accept **Conventional Path**, and then click **Next**.
 - h) On the Load Data: Options page, accept all defaults, but enter **/home/oracle/labs/lab_18_05.log** as log file name and path.
 - i) Review the advanced options, but do not change any, and then click **Next**.

Solutions for Practice 18: Moving Data (continued)

- j) On the Load Data: Schedule page, enter **lab_18_05** as **Job Name** and **Load data into the PRODUCT_MASTER table** as **Description**. Let the job start immediately, and then click **Next**.

Control File Data File Load Method Options **Schedule** Review

Load Data: Schedule

Database **orcl** Cancel Back **Step 5 of 6** Next

Specify a name and description for the load data job. Specify a date to start the job.

Job Parameters

Job Name **lab_18_05**

Description **Load data into PRODUCT_MASTER table**

Job Schedule

Start

☒ Immediately
☐ Later

Date **Aug 2, 2005**

(example: Aug 2, 2005)

Time **5** **15** ☐ AM ☒ PM

- k) On the Load Data: Review page, review the loading information and parameters, and then click **Submit Job**.
- l) Confirm your results by viewing your `lab_18_05.log` file in your `/home/oracle/labs` directory.
6. As the INVENTORY user, load data into the PRODUCT_ON_HAND table by using SQL*Loader command line.
- a) Invoke a terminal window and navigate to the `/home/oracle/labs` directory.
- b) Enter the following SQL*Loader command (in continuation, without pressing [Enter] before reaching the end of the command): **sqlldr**
userid=inventory/verysecure control=lab_18_06.ctl log=lab_18_06.log data=lab_18_06.dat.

```
sqlldr userid=inventory/verysecure control=lab_18_06.ctl
log=lab_18_06.log data=lab_18_06.dat

SQL*Loader: Release 10.2.0.1.0 - Production on Tue Aug 02 22:24:44
2005

Copyright © 1982, 2005, Oracle. All rights reserved.

Commit point reached - logical record count 64
Commit point reached - logical record count 82
Commit point reached - logical record count 83
```

Solutions for Practice 18: Moving Data (continued)

- c) Confirm your results by viewing your `lab_18_06.log` file in your `/home/oracle/labs` directory.

Congratulations!

*You completed all practices for the
Oracle Database 10g: Administration Workshop I*

